

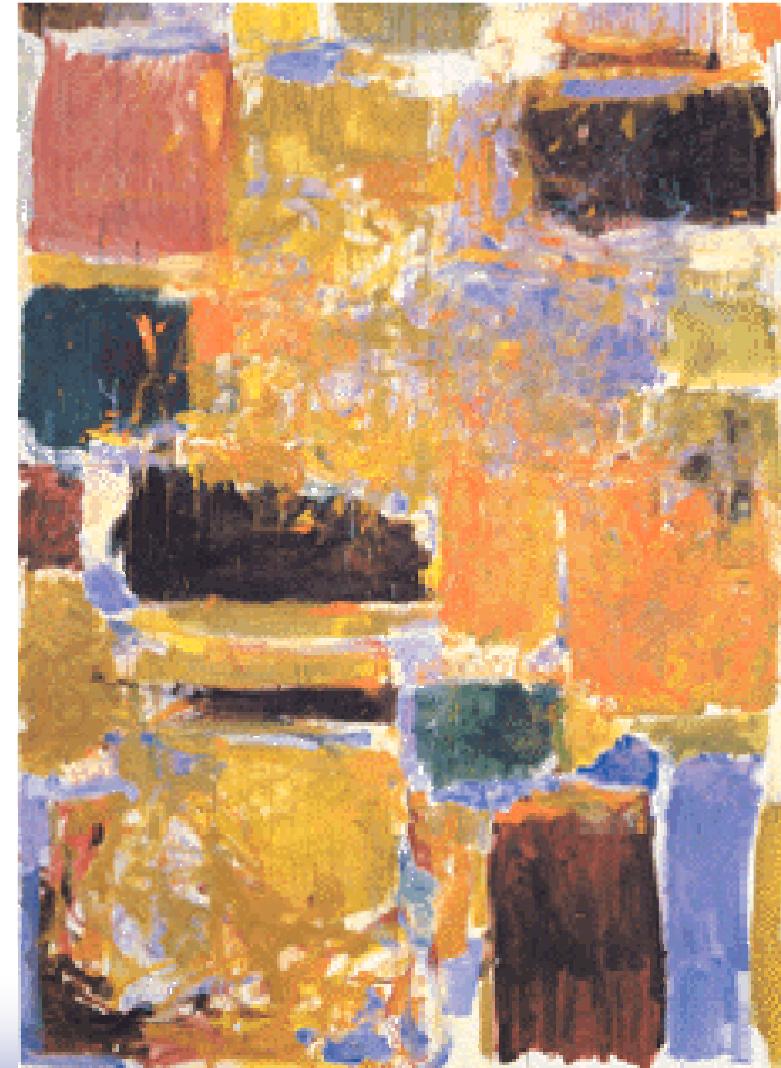


CS150 - EE141/241A
Fall 2014
Digital Design and
Integrated Circuits

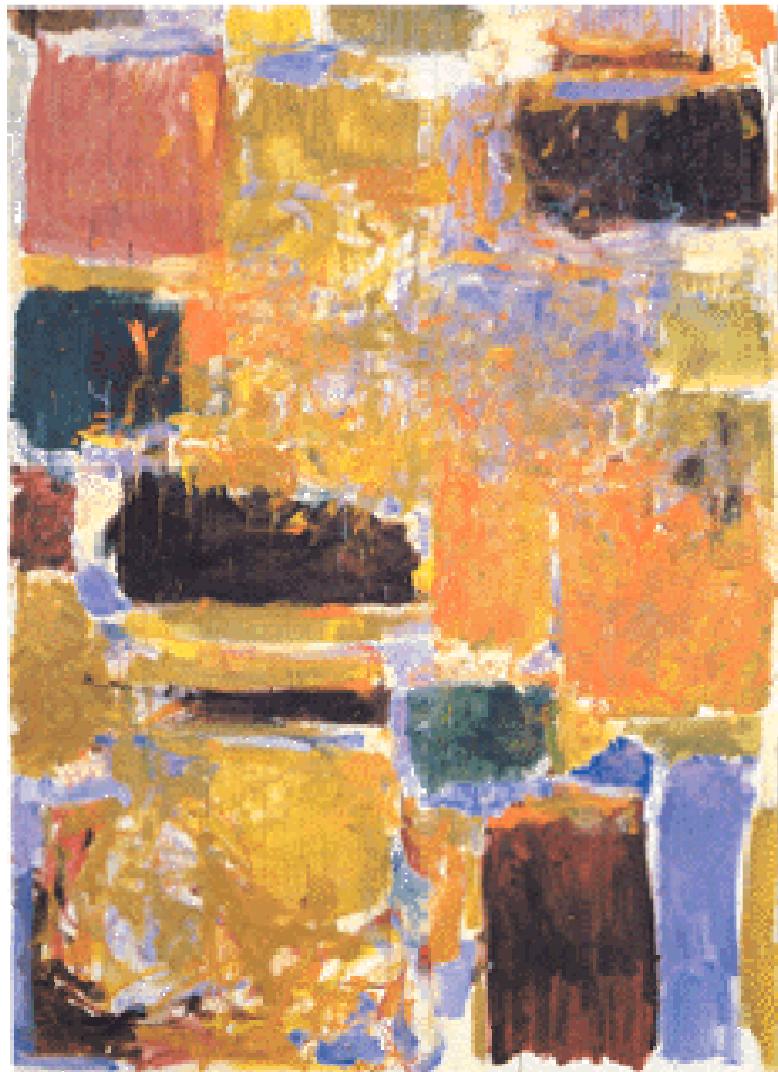
Instructors:
John Wawrzynek and Vladimir Stojanovic

Lecture 22

Outline

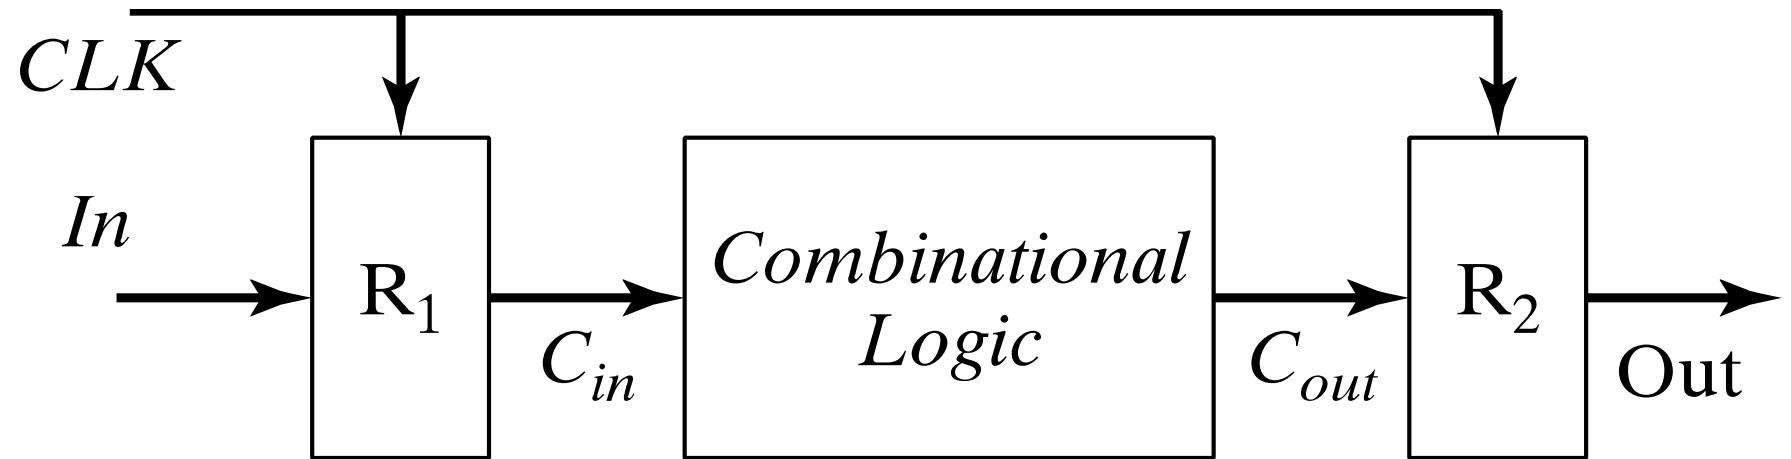


- Clock nonidealities
- Timing
- Clock Distribution

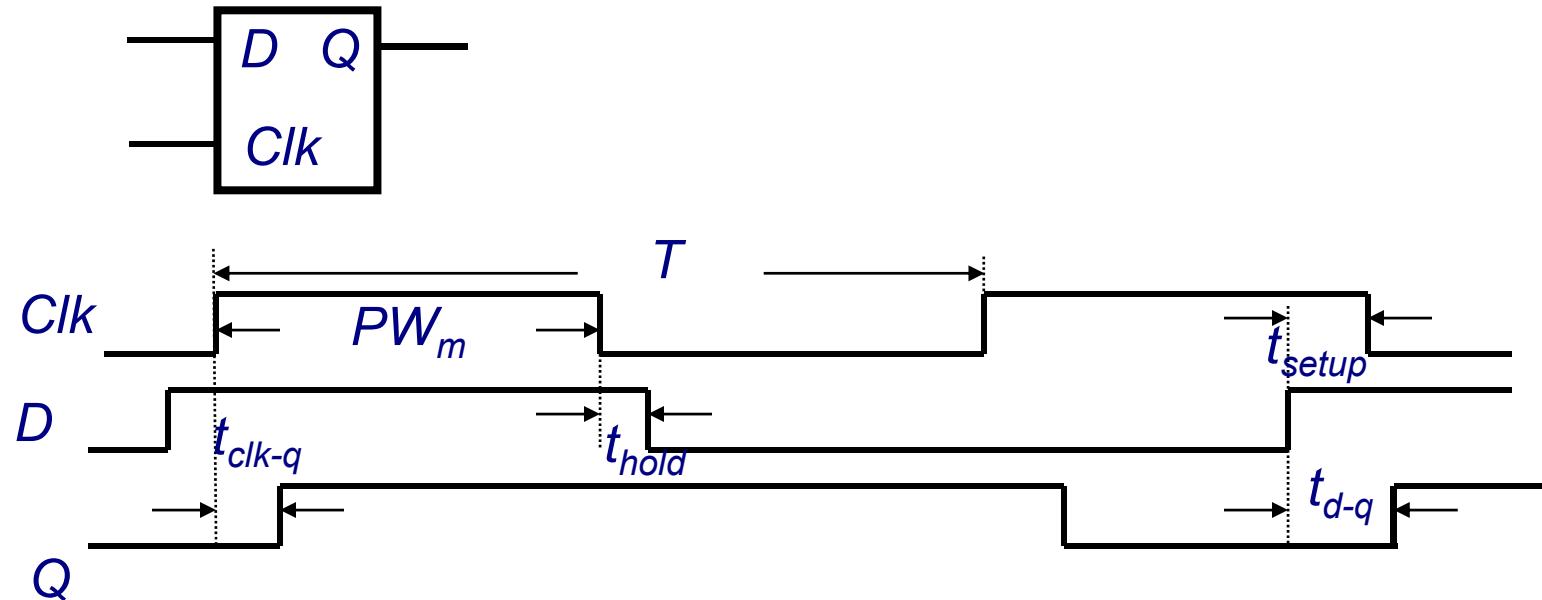


Synchronous Timing - Review

Synchronous Timing

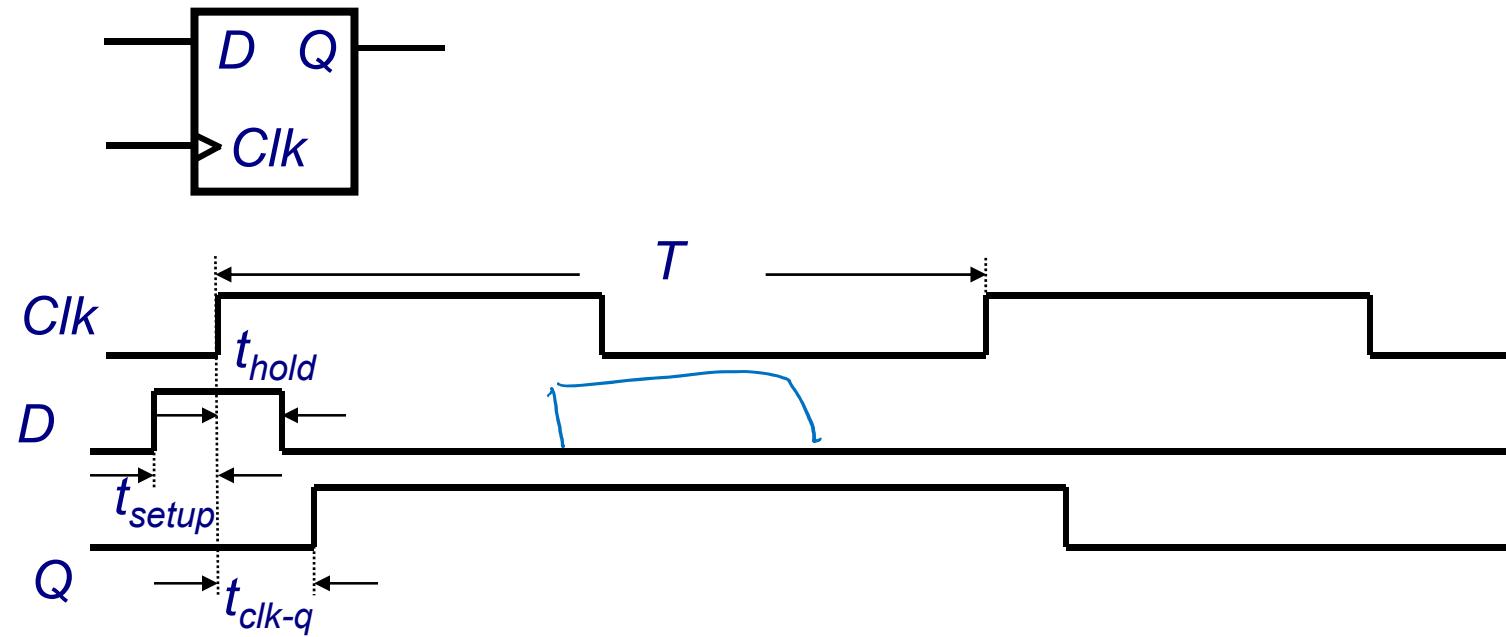


Latch Parameters



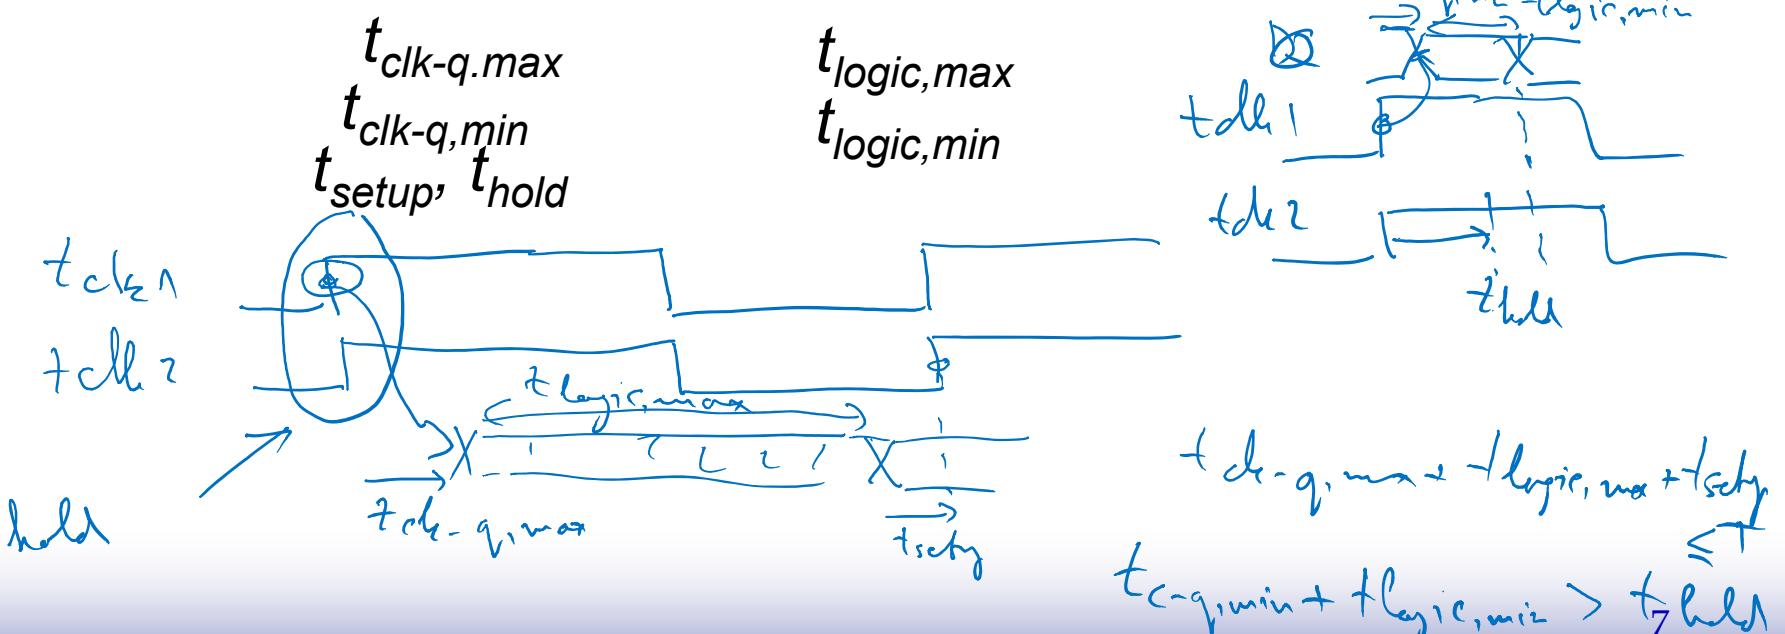
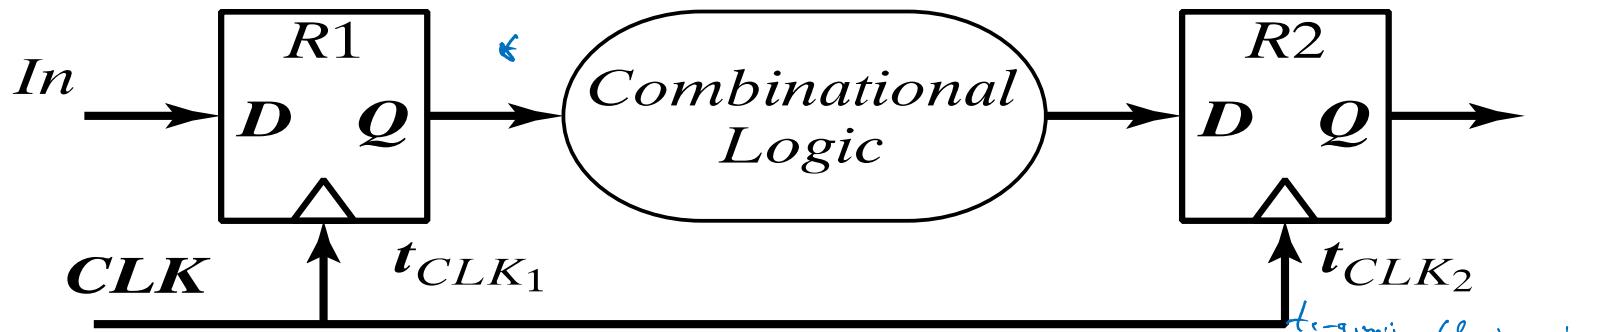
Delays can be different for rising and falling data transitions

Register Parameters

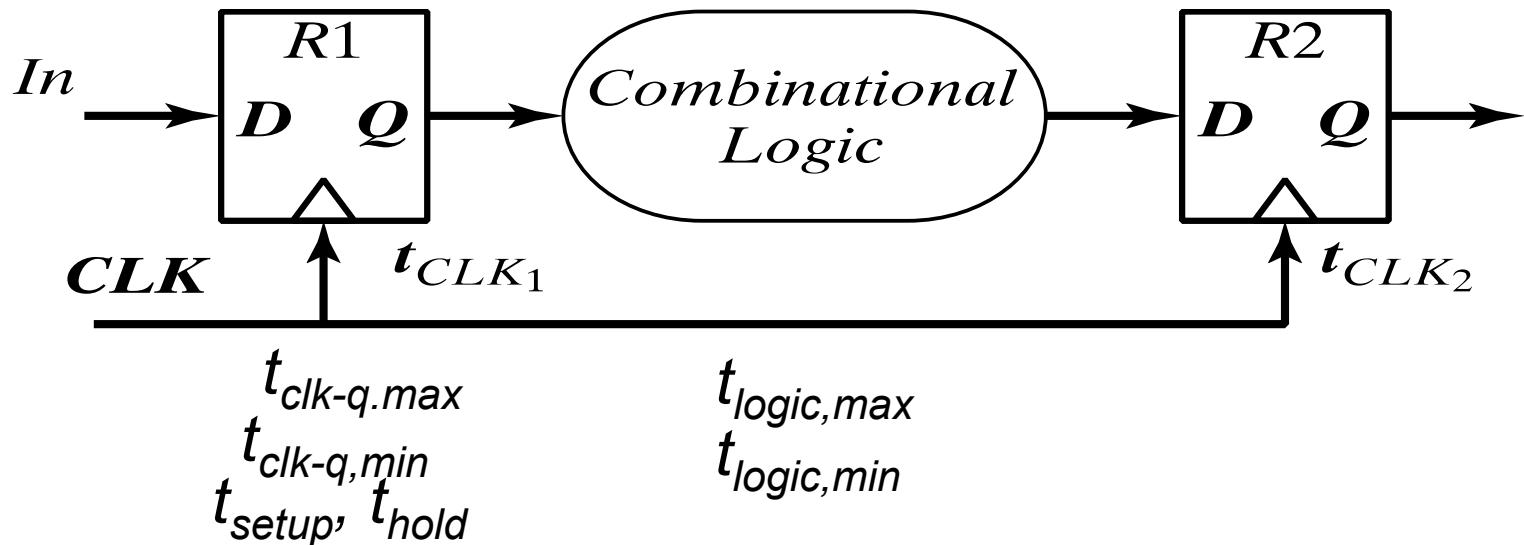


Delays can be different for rising and falling data transitions

Timing Constraints

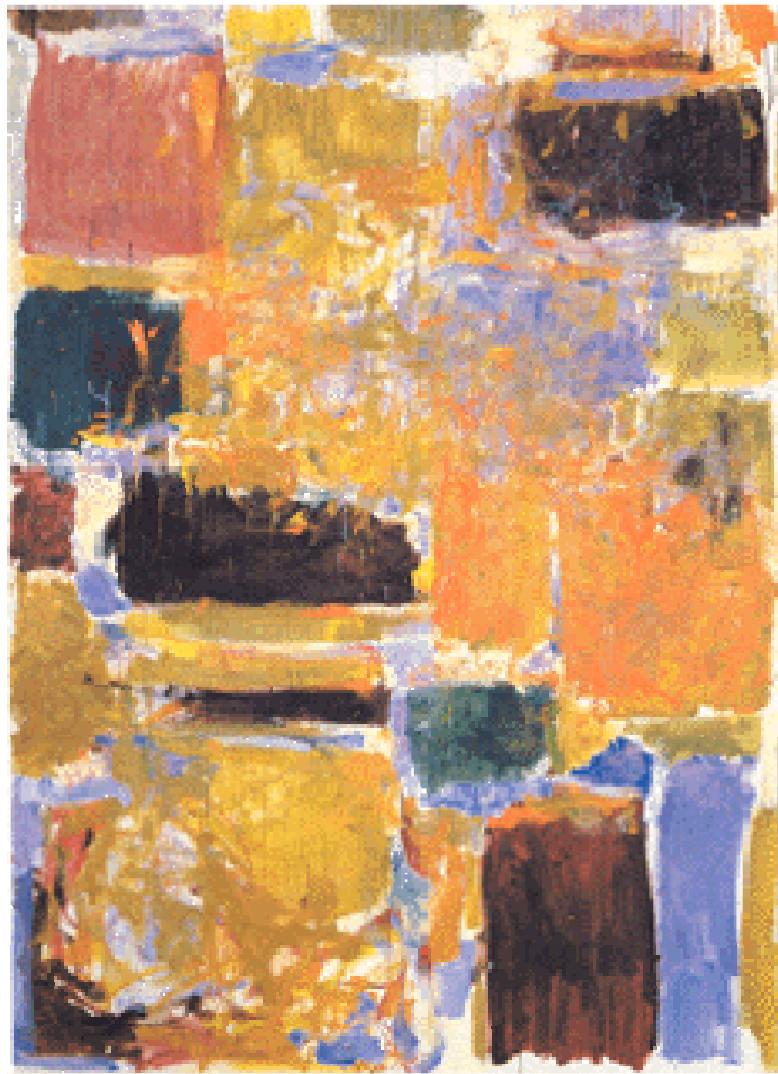


Timing Constraints



Cycle time (max): $T_{Clk} > t_{clk-q,max} + t_{logic,max} + t_{setup}$

Race margin (min): $t_{hold} < t_{clk-q,min} + t_{logic,min}$

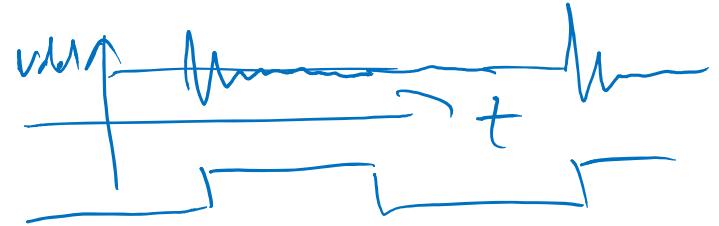


Clock Nonidealities

Clock Nonidealities

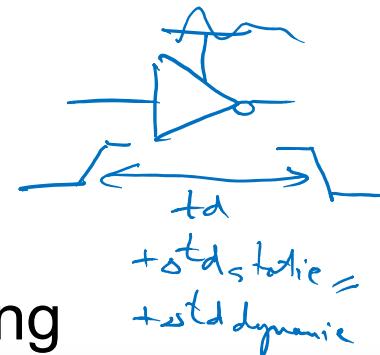
□ Clock skew

- Time difference between the sink (receiving) and source (launching) edge
- Spatial variation in temporally equivalent clock edges; deterministic + random, t_{SK}



□ Clock jitter

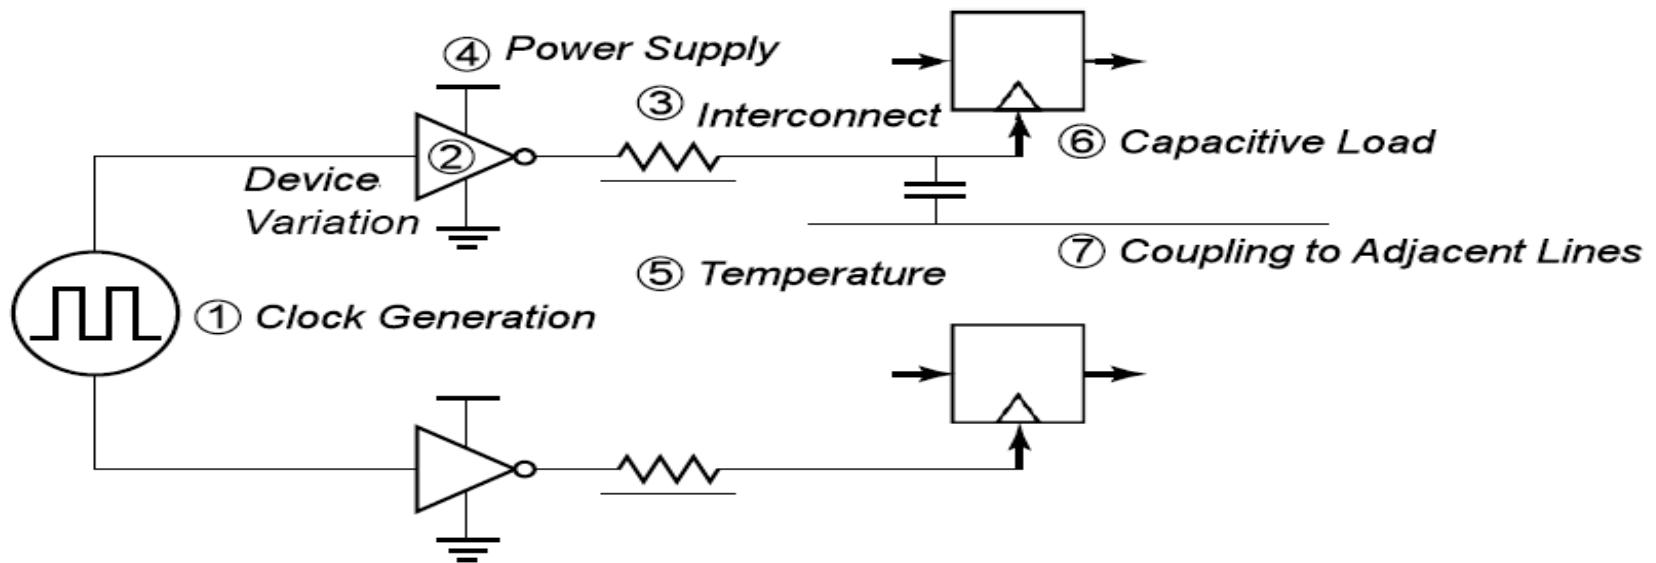
- Temporal variations in consecutive edges of the clock signal; modulation + random noise
- Cycle-to-cycle (short-term) t_{JS}
- Long term t_{JL}



□ Variation of the pulse width

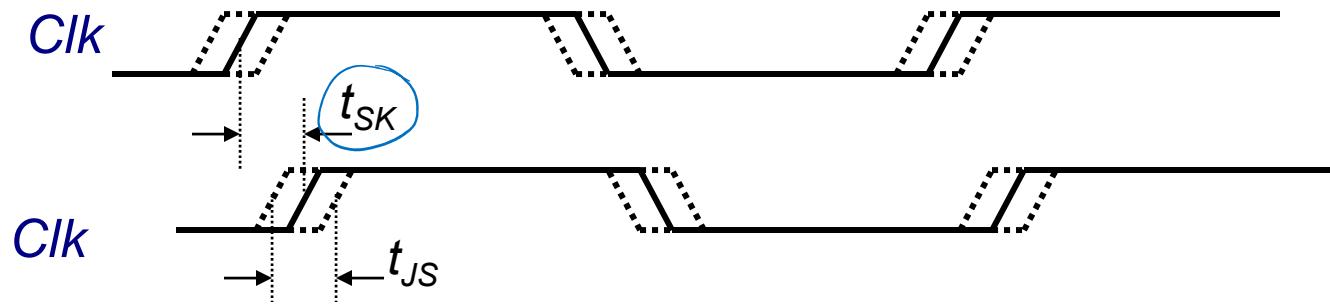
- Important for level sensitive clocking

Clock Uncertainties

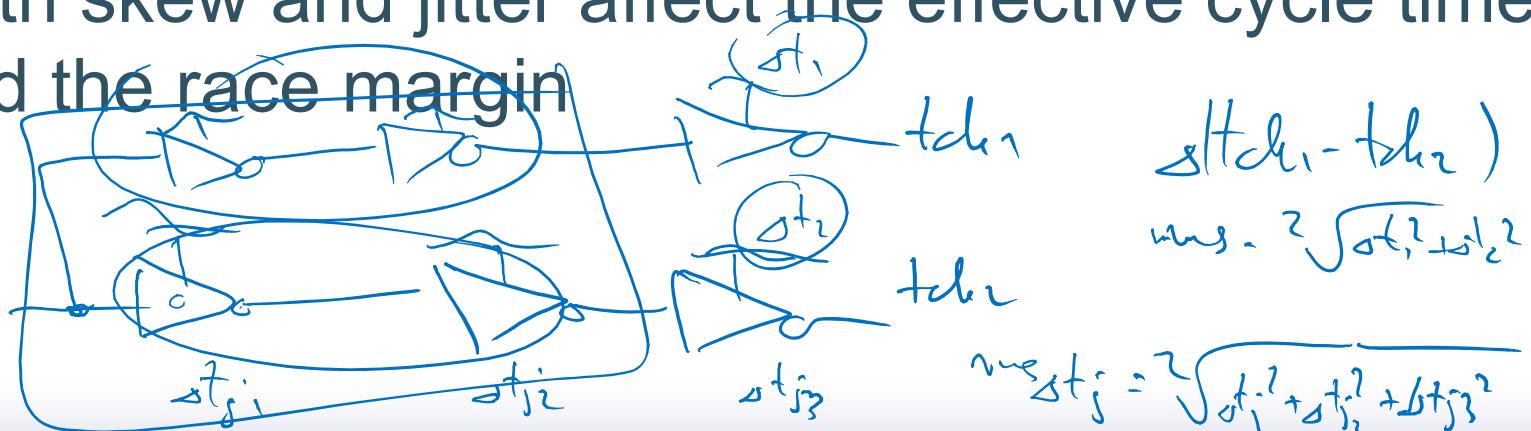


Sources of clock uncertainty

Clock Skew and Jitter

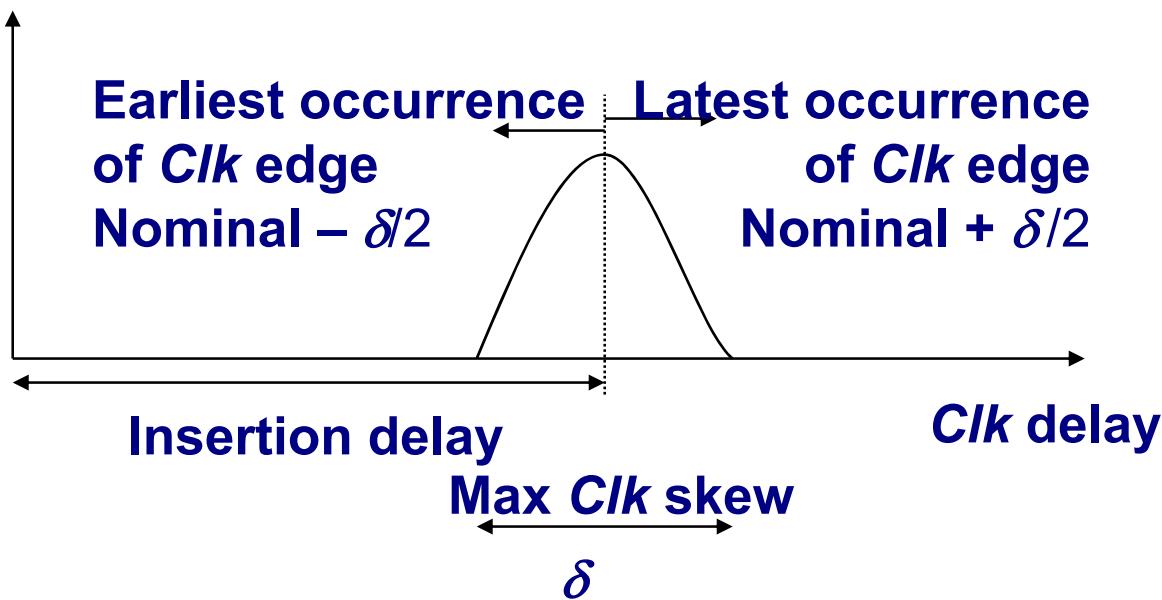


- Both skew and jitter affect the effective cycle time and the race margin

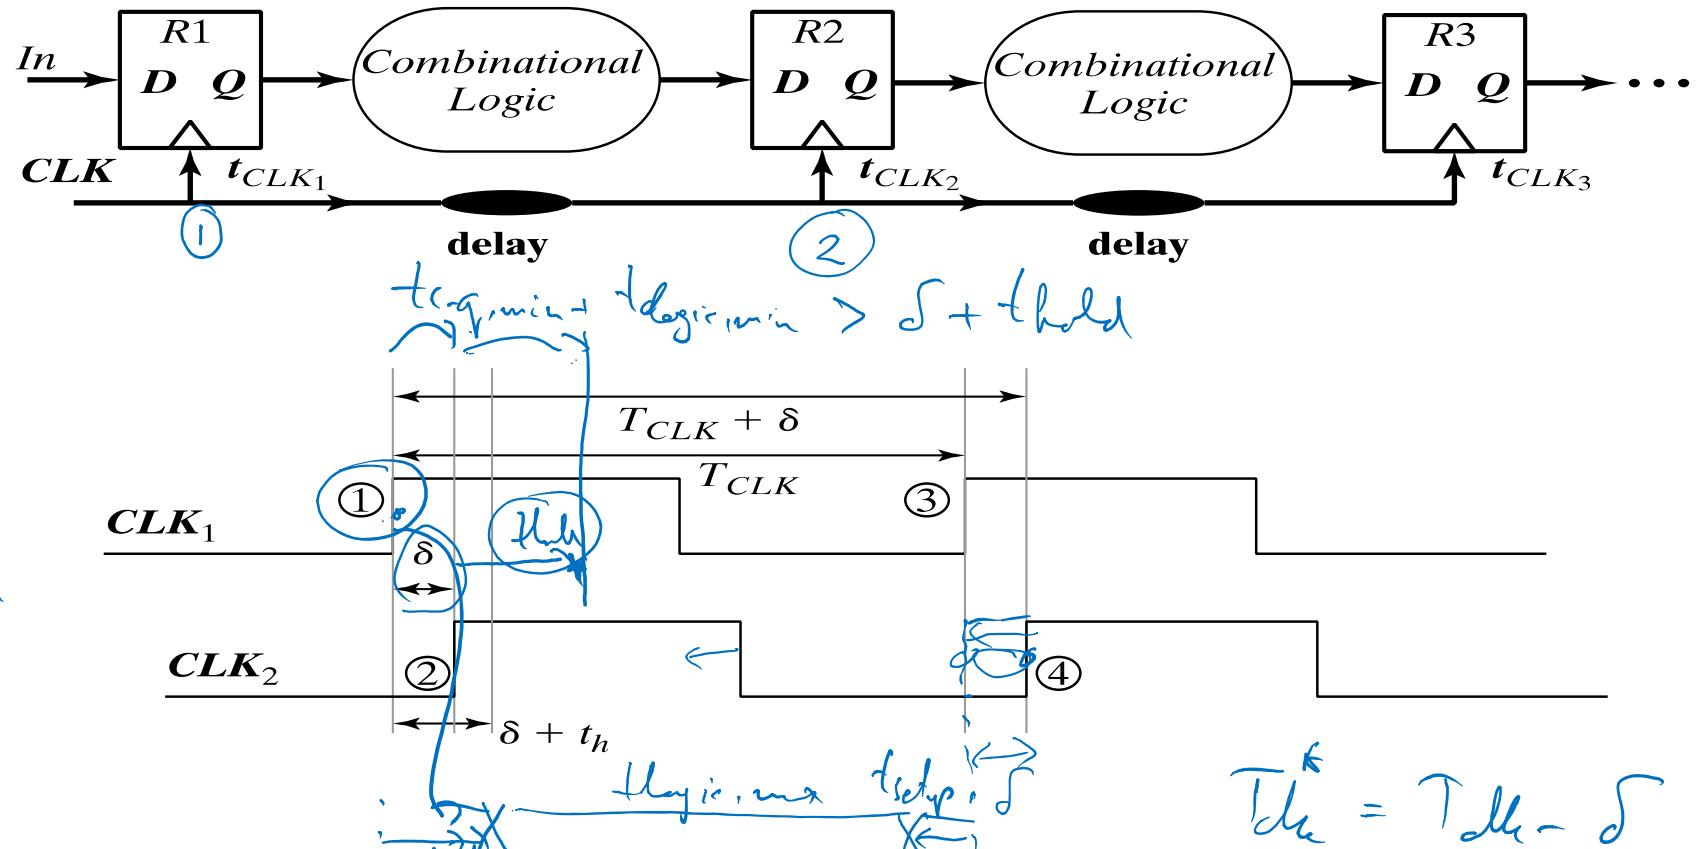


Clock Skew

of registers

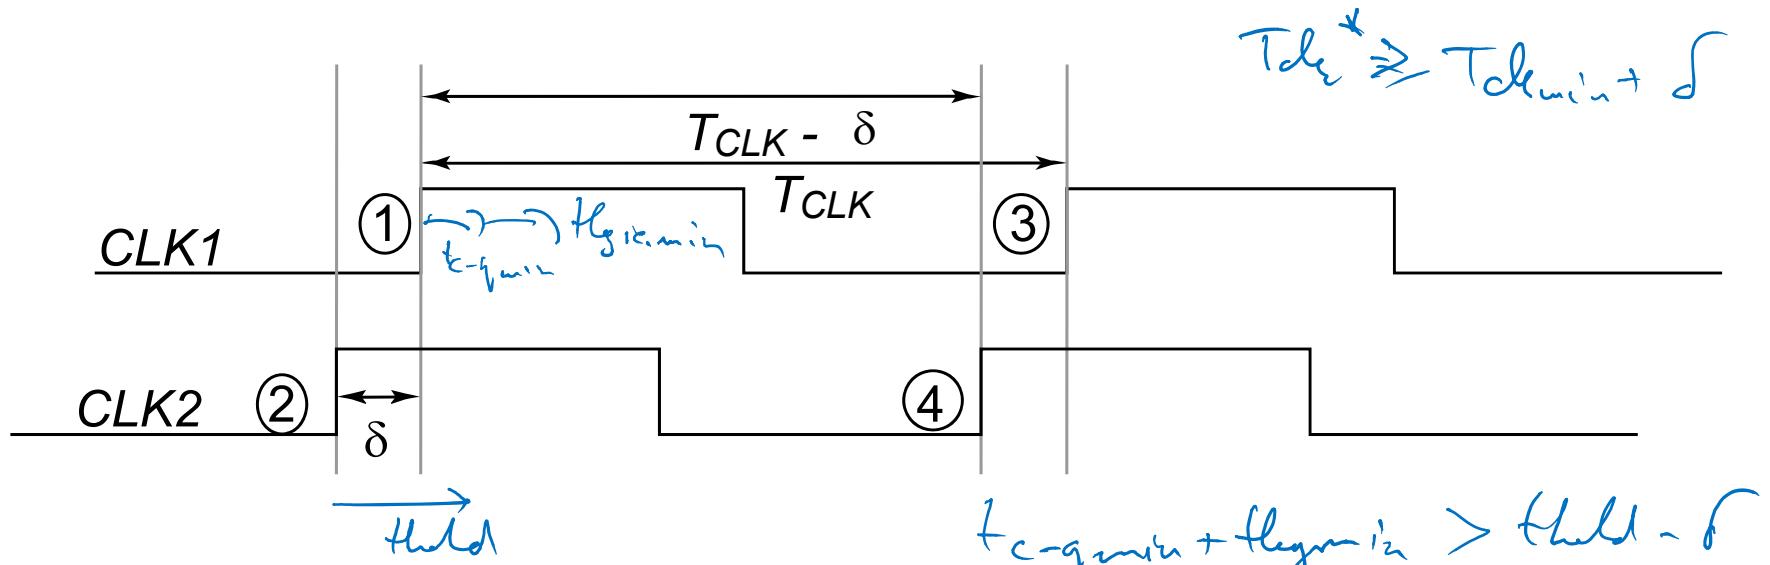
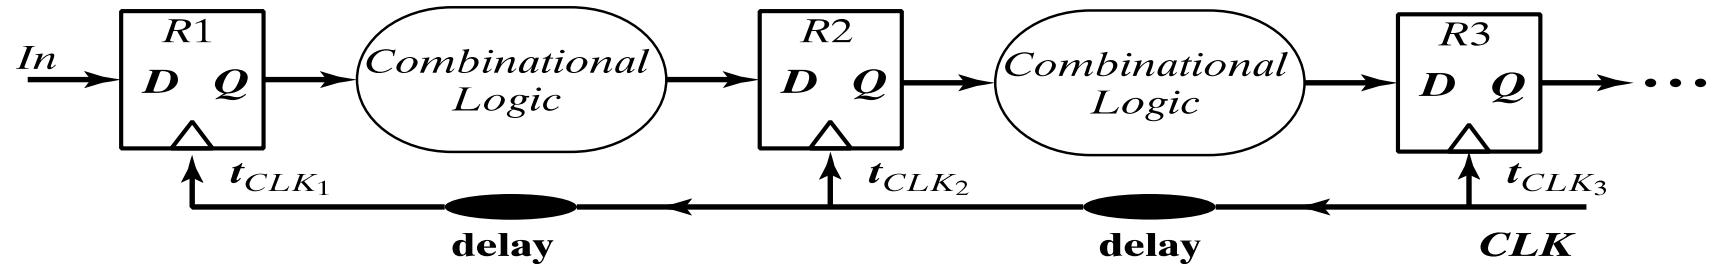


Positive Skew

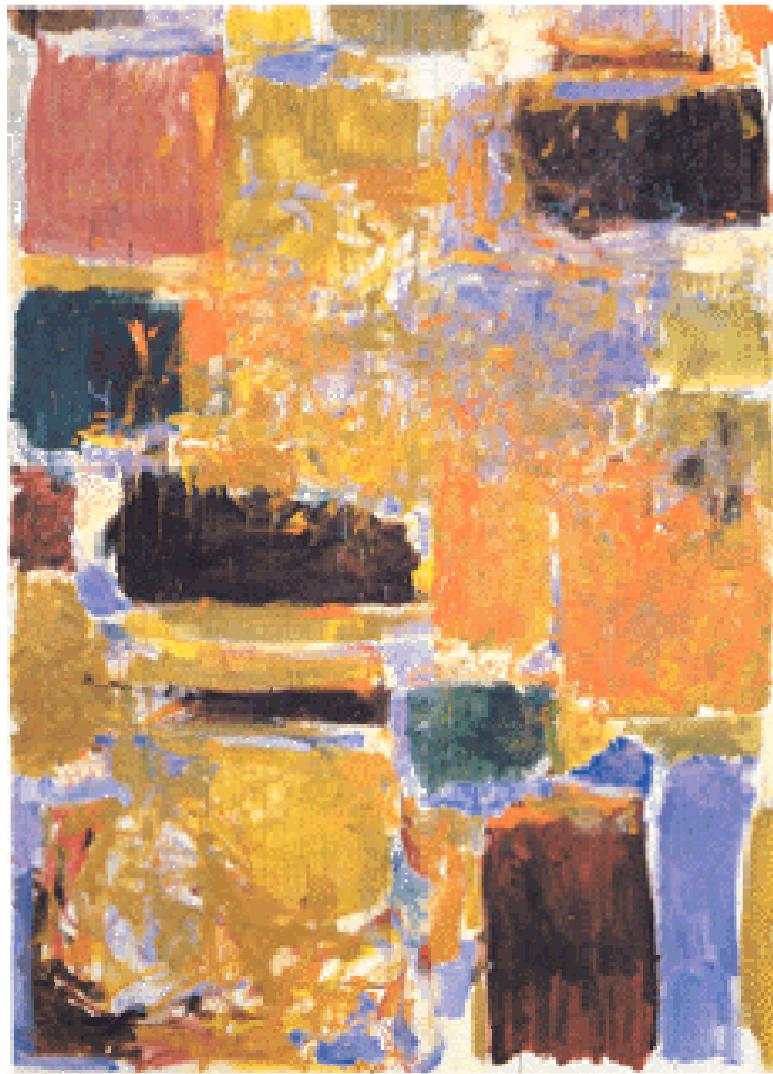


Launching edge arrives before the receiving edge

Negative Skew

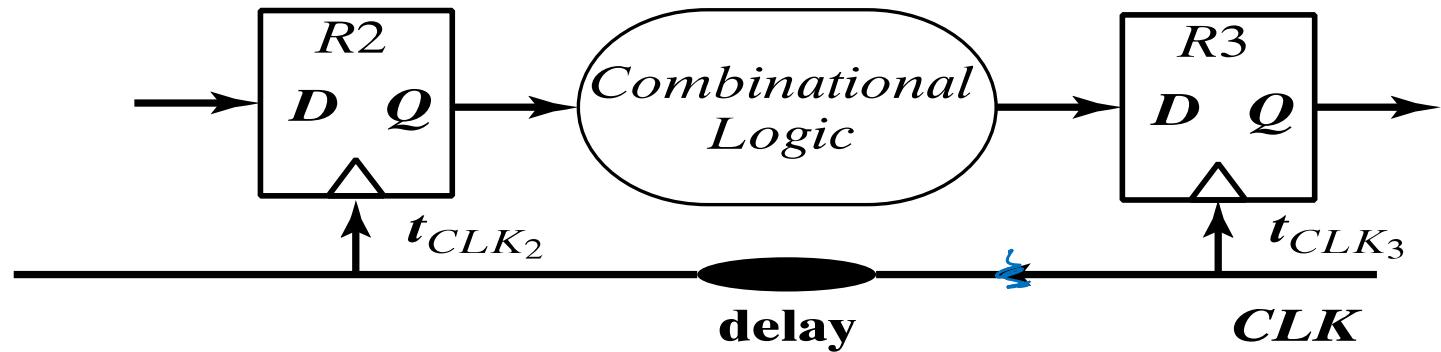


Receiving edge arrives before the launching edge



Timing with Clock Nonidealities

Timing Constraints



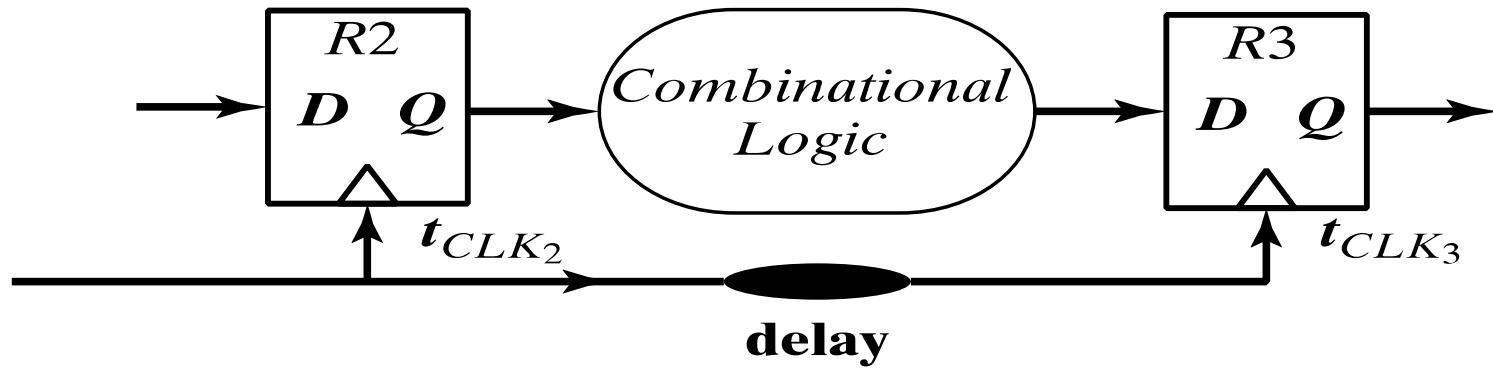
$t_{clk-q,max}$
 $t_{clk-q,min}$
 t_{setup}, t_{hold}

$t_{logic,max}$
 $t_{logic,min}$

Minimum cycle time:

$$T_{clk} + \delta = t_{clk-q,max} + t_{setup} + t_{logic,max}$$

Timing Constraints

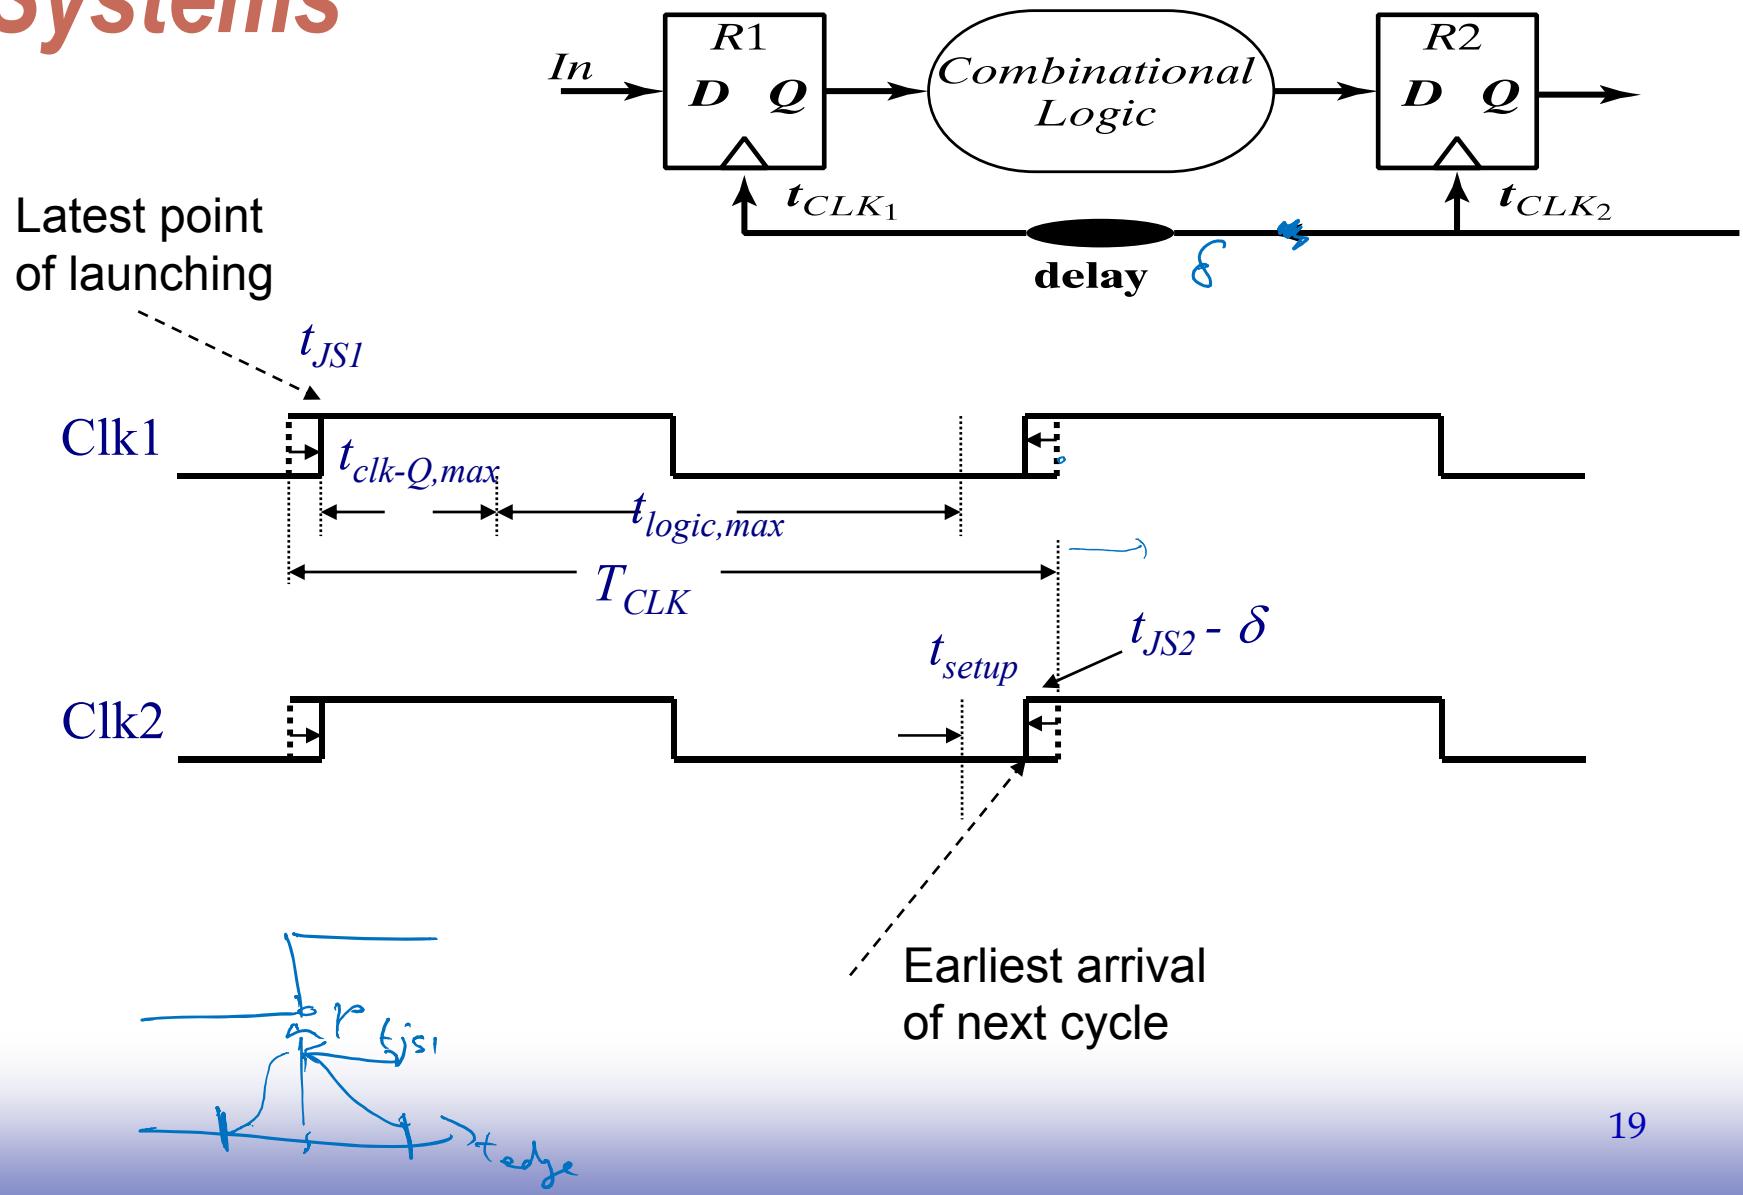


$$\begin{array}{ll} t_{clk-q,max} & t_{logic,max} \\ t_{clk-q,min} & t_{logic,min} \\ t_{setup}, t_{hold} & \end{array}$$

Hold time constraint:

$$t_{(clk-q,min)} + t_{(logic,min)} > t_{hold} + \delta$$

Longest Logic Path in Edge-Triggered Systems



Clock Constraints in Edge-Triggered Systems

If launching edge is late and receiving edge is early, the data will not be too late if:

$$t_{clk-q,max} + t_{logic,max} + t_{setup} < T_{CLK} - t_{JS,1} - t_{JS,2} + \delta$$

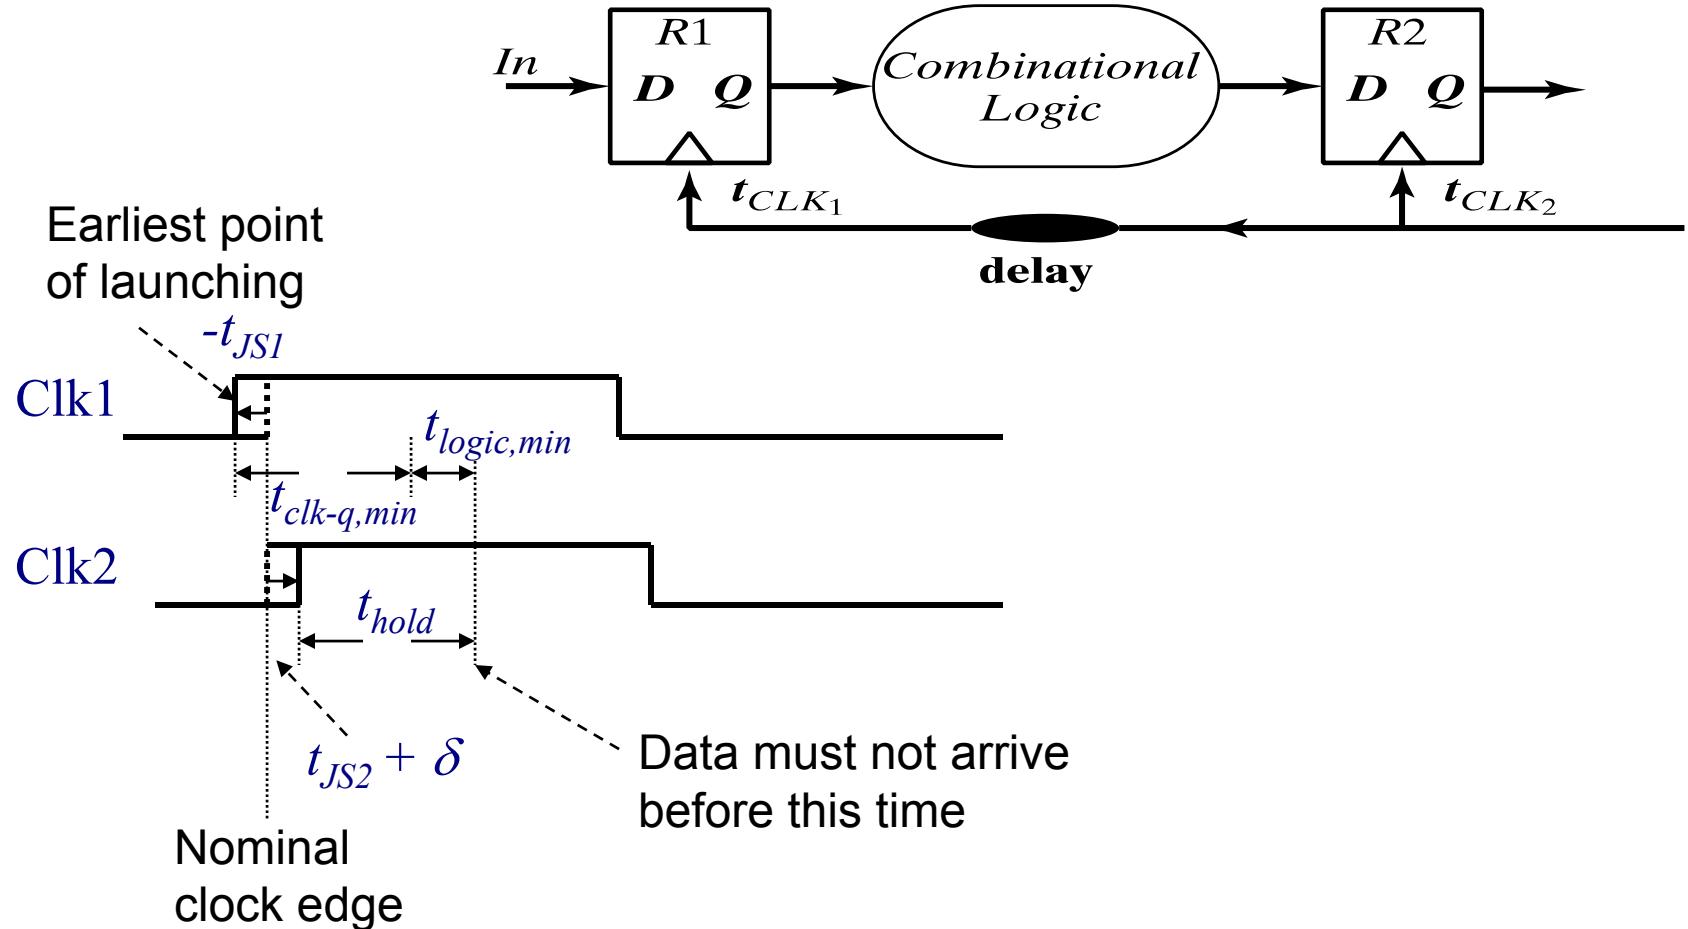
Minimum cycle time is determined by the maximum delays through the logic

$$t_{clk-q,max} + t_{logic,max} + t_{setup} - \delta + \underbrace{2t_{JS}}_{\text{peak-to-peak jitter}} < T_{CLK}$$

Skew can be either positive or negative

Jitter t_{JS} usually expressed as peak-to-peak or $N \times \text{rms}$ value

Shortest Path



Clock Constraints in Edge-Triggered Systems

If launching edge is early and receiving edge is late:

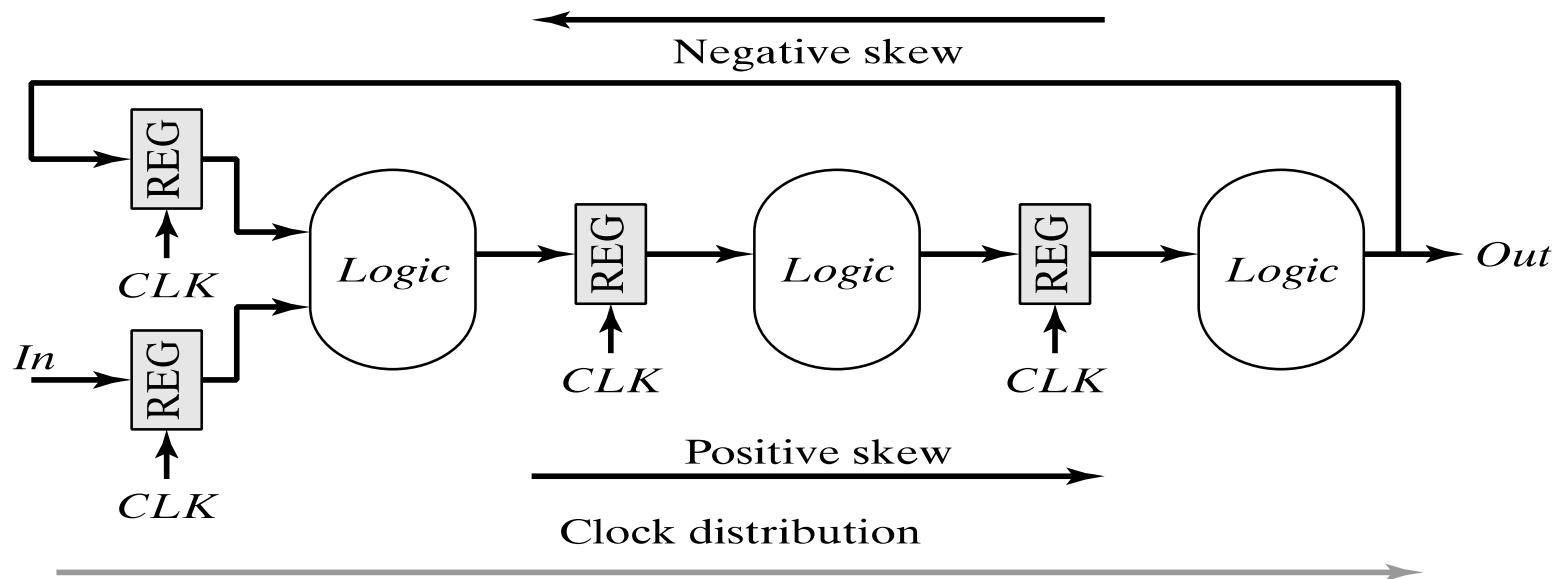
$$t_{clk-q,min} + t_{logic,min} - t_{JS,1} > t_{hold} + t_{JS,2} + \delta$$

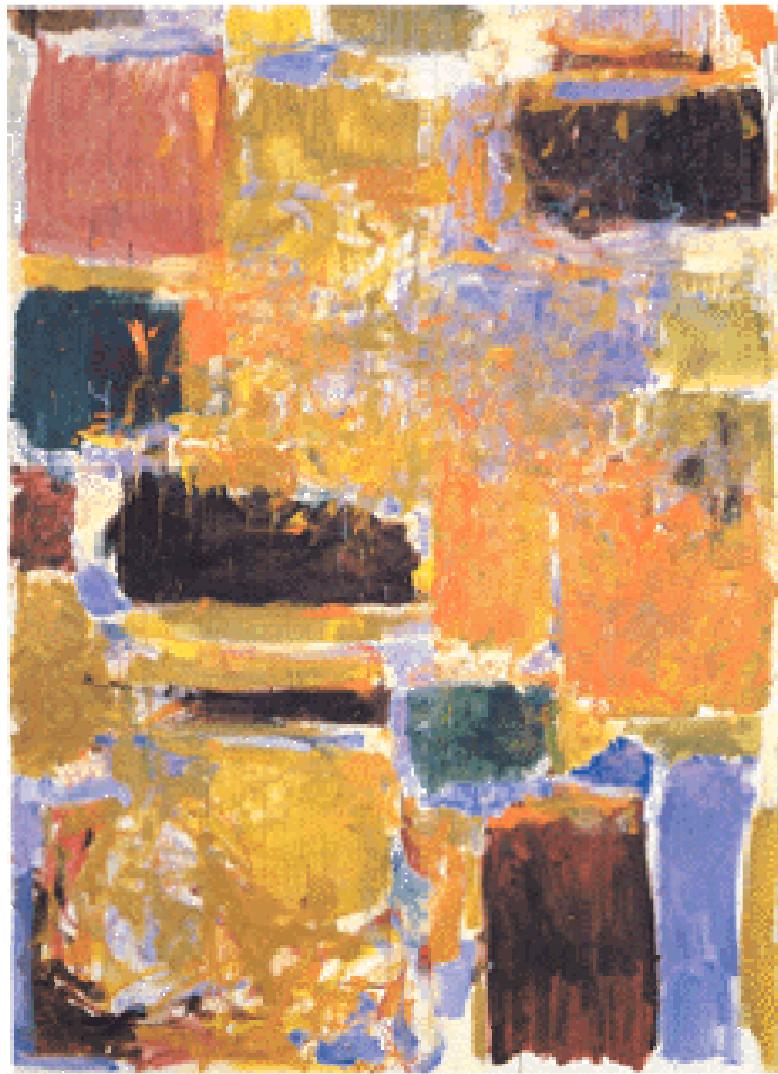
Minimum logic delay

$$t_{clk-q,min} + t_{logic,min} > t_{hold} + 2t_{JS} + \delta$$

(This assumes jitter at launching and receiving clocks are independent – which usually is not true)

Datapath with Feedback





Clock Distribution

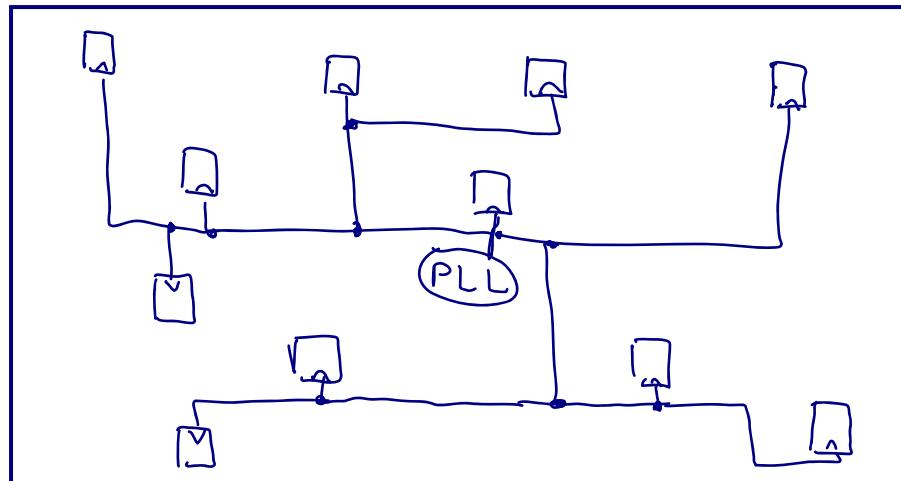
1

Clock Distribution

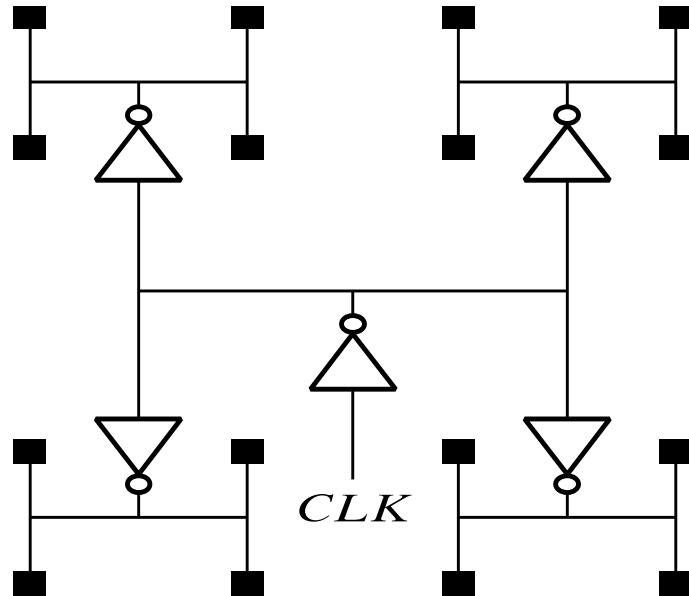
- Single clock generally used to synchronize all logic on the same chip
 - Need to distribute clock over the entire die
 - While maintaining low skew/jitter
 - (And without burning too much power)

Clock Distribution

- ❑ What's wrong with just routing wires to every point that needs a clock?

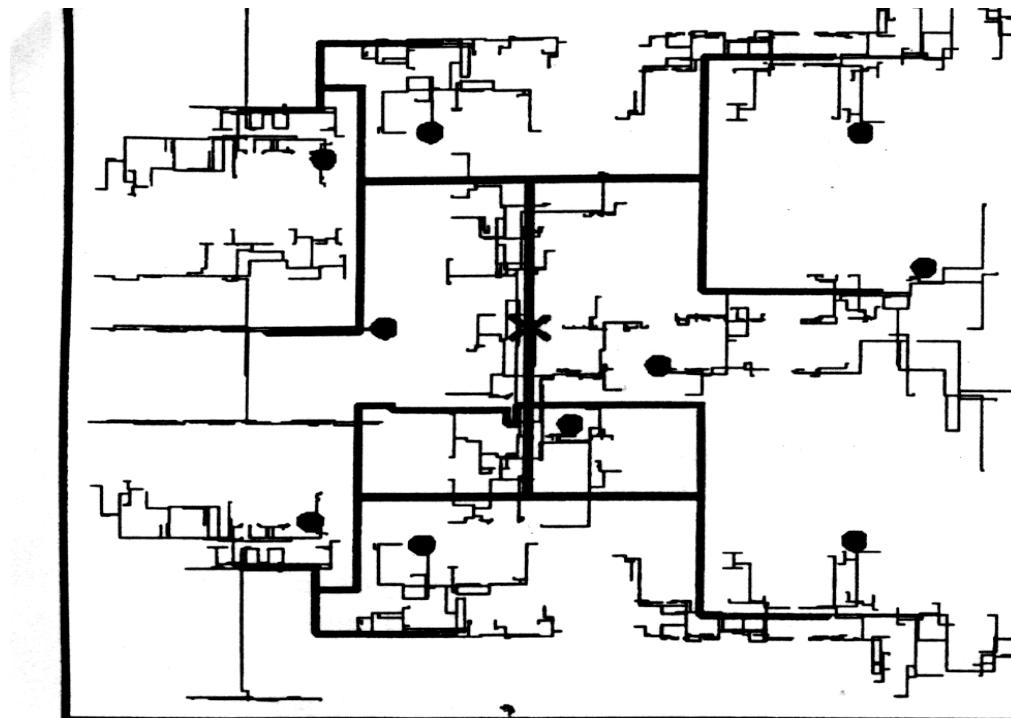


H-Tree



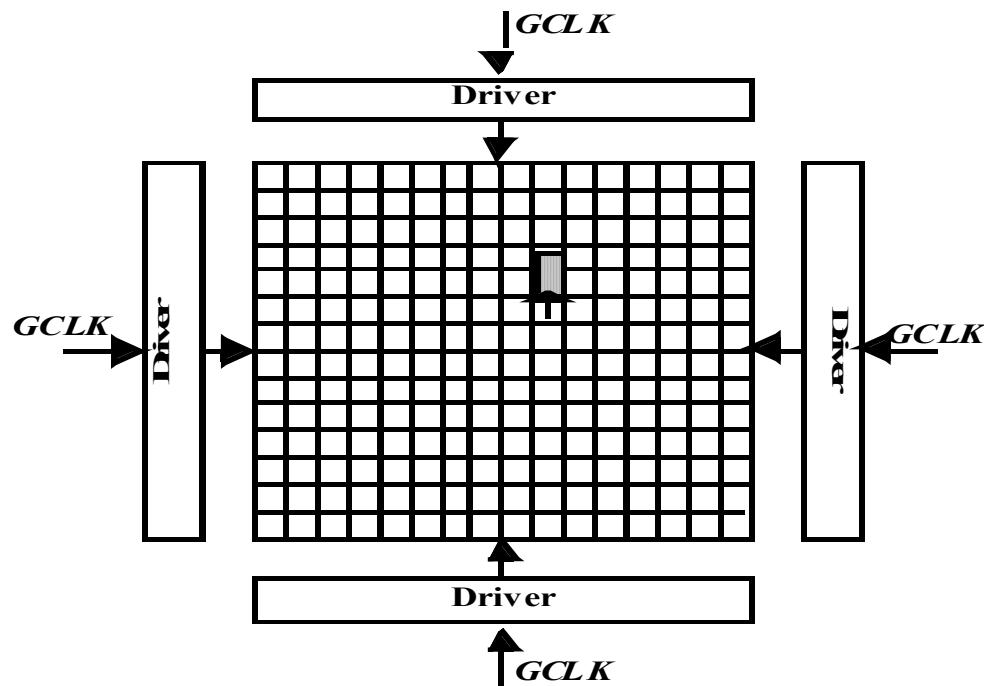
Equal wire length/number of buffers to get to every location

More realistic H-tree



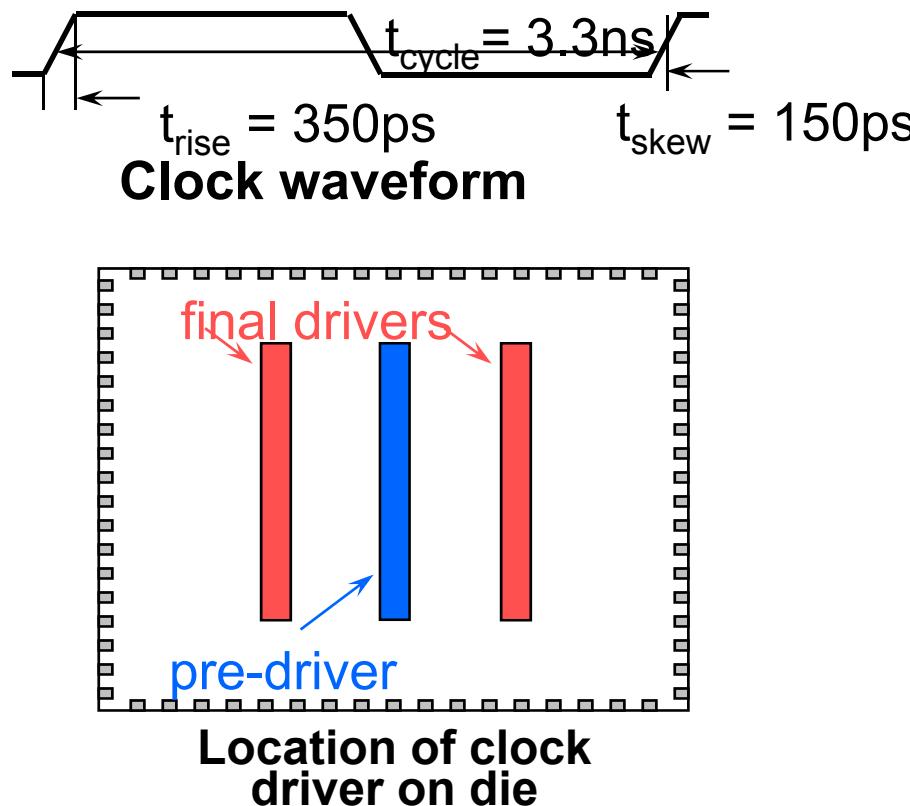
[Restle98]

Clock Grid

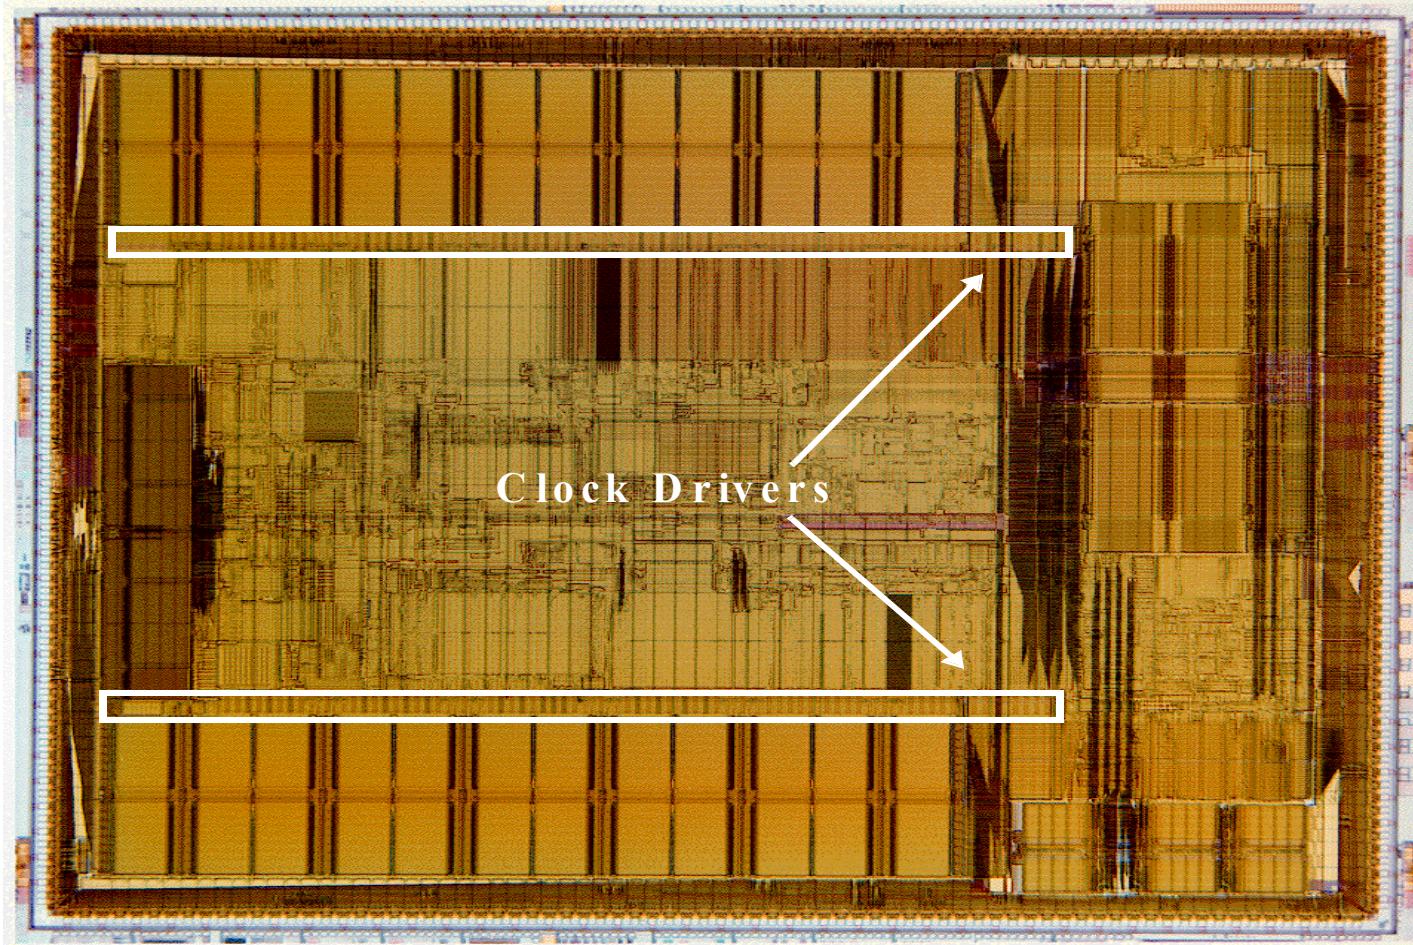


- No RC matching
- But huge power

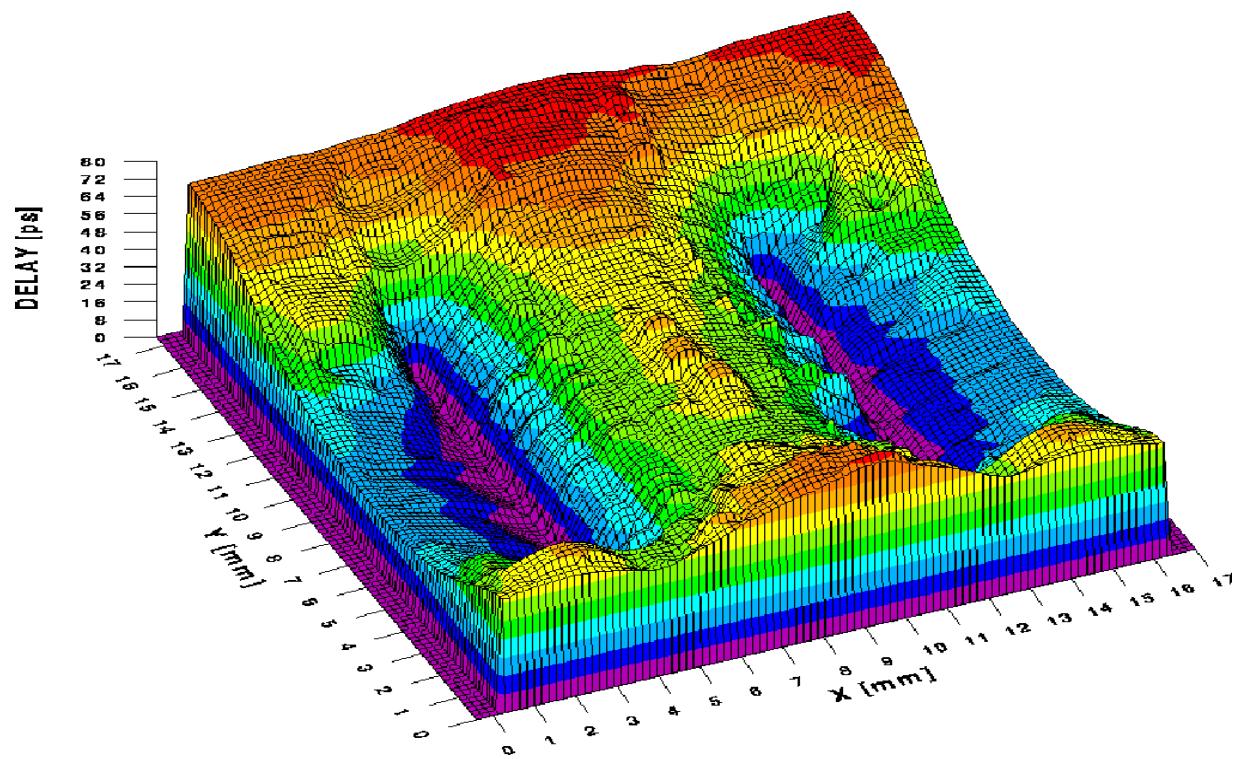
Example: DEC Alpha 21164 (1995)



- ❑ 2 phase single wire clock, distributed globally
- ❑ 2 distributed driver channels
 - Reduced RC delay/skew
 - Improved thermal distribution
 - 3.75nF clock load, 20W power
 - 58 cm final driver width
- ❑ Local inverters for latching
- ❑ Conditional clocks in caches to reduce power
- ❑ More complex race checking
- ❑ Device variation

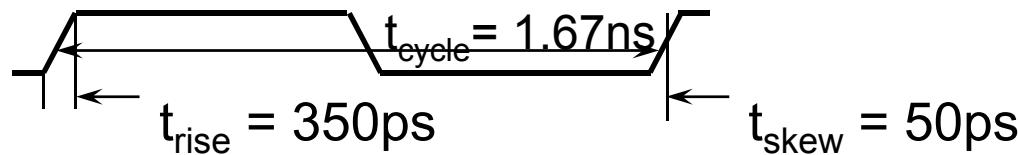


Clock Skew in Alpha Processor

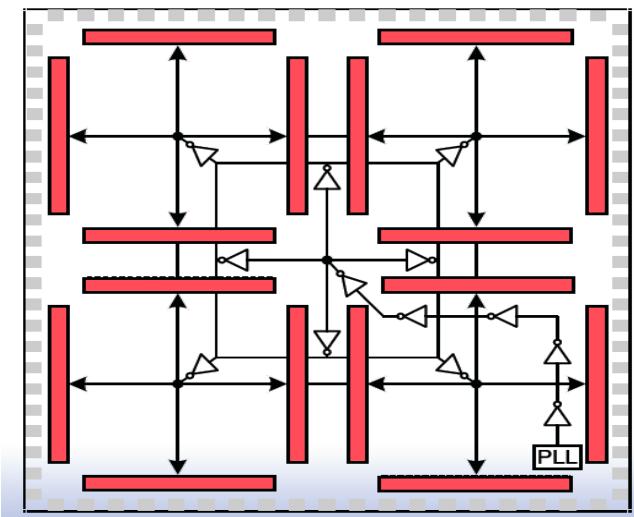


EV6 (Alpha 21264) Clocking

600 MHz – 0.35 micron CMOS

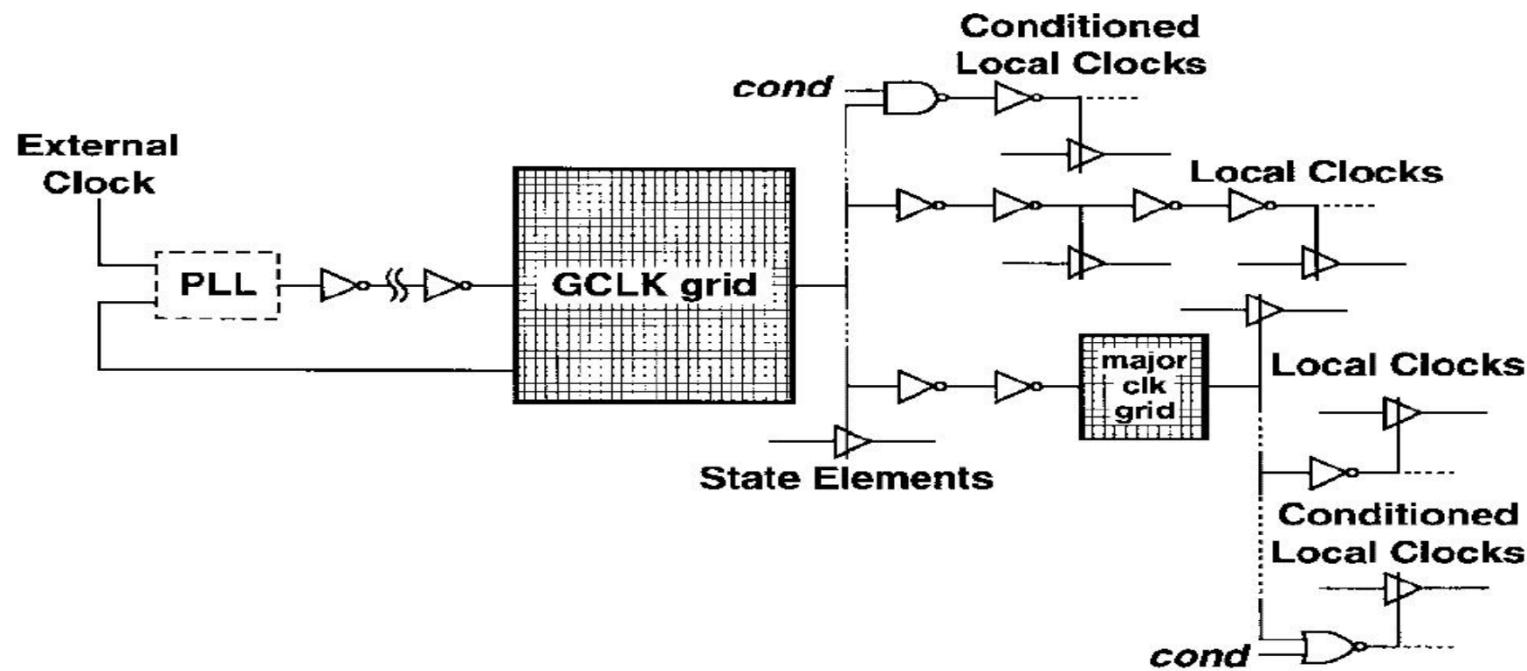


Global clock waveform

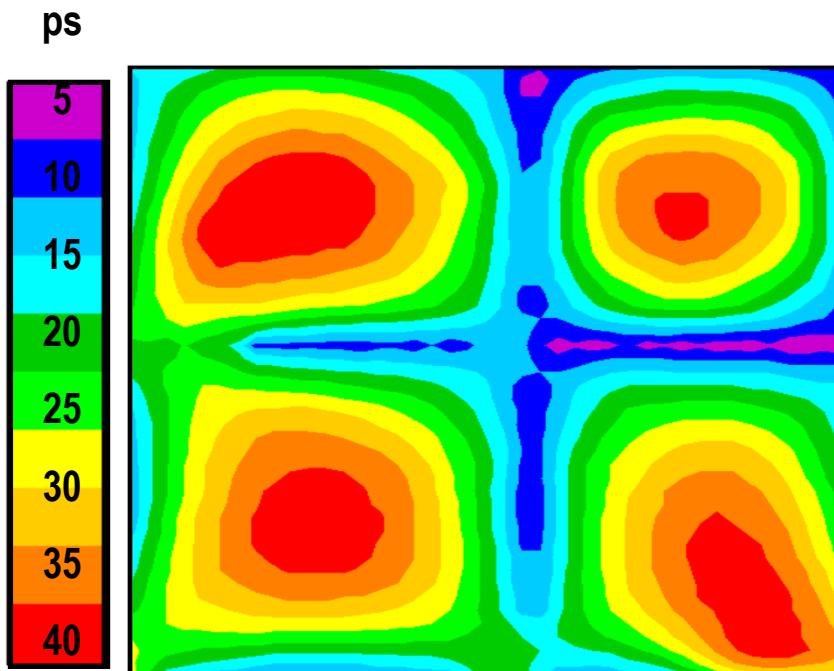


- 2 Phase, with multiple conditional buffered clocks
 - 2.8 nF clock load
 - 40 cm final driver width
- Local clocks can be gated “off” to save power
- Reduced load/skew
- Reduced thermal issues
- Multiple clocks complicate race checking

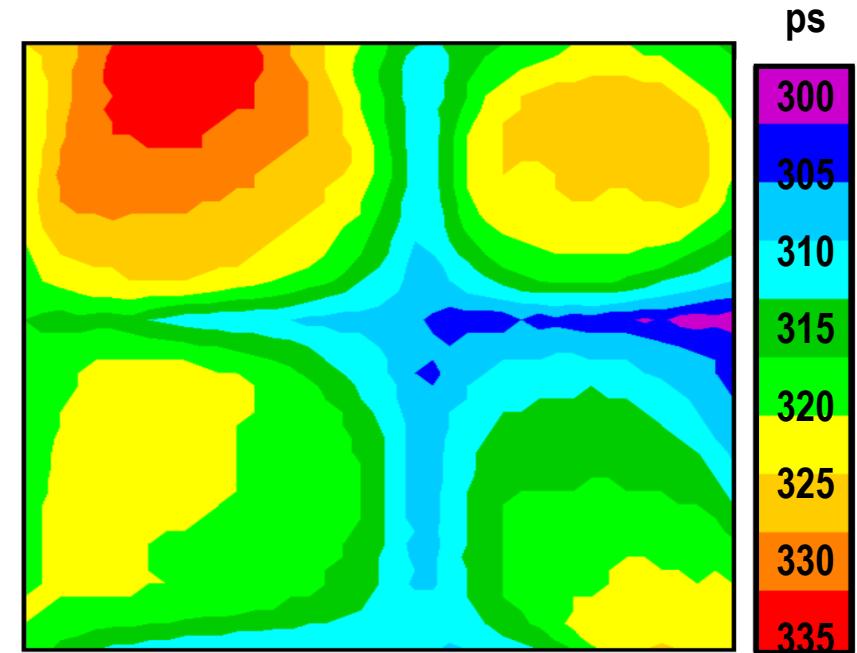
21264 Clocking



EV6 Clock Results



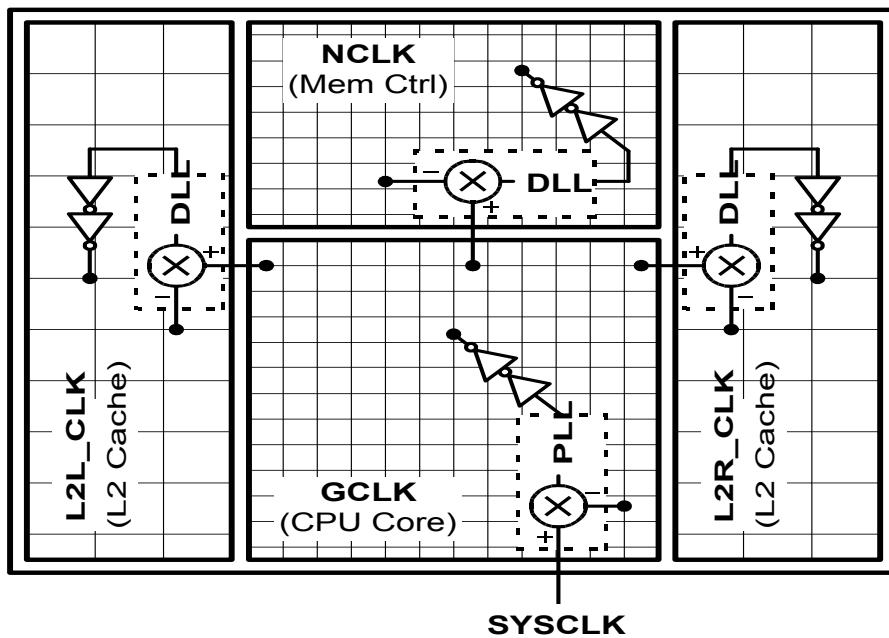
45
GCLK Skew
50 (at Vdd/2 Crossings)



340
GCLK Rise Times
(20% to 80% Extrapolated to 0% to 100%)

EV7 Clock Hierarchy (2002)

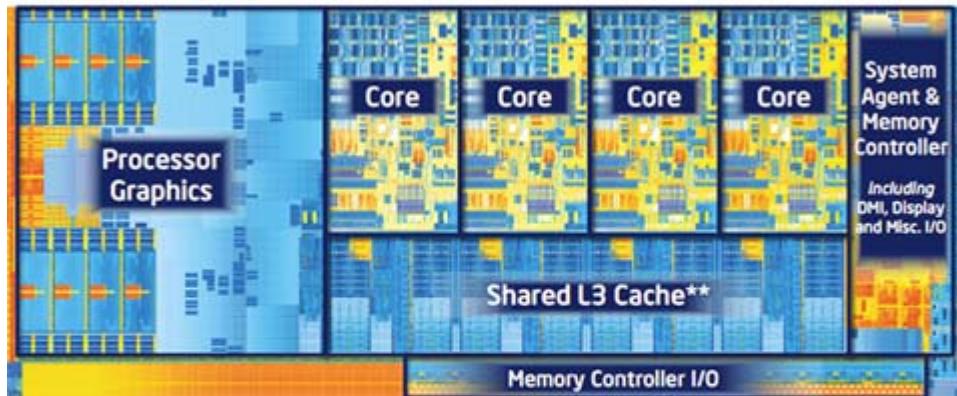
Active Skew Management and Multiple Clock Domains



- + widely dispersed drivers
- + DLLs compensate static and low-frequency variation
- + divides design and verification effort
- DLL design and verification is added work
- + tailored clocks

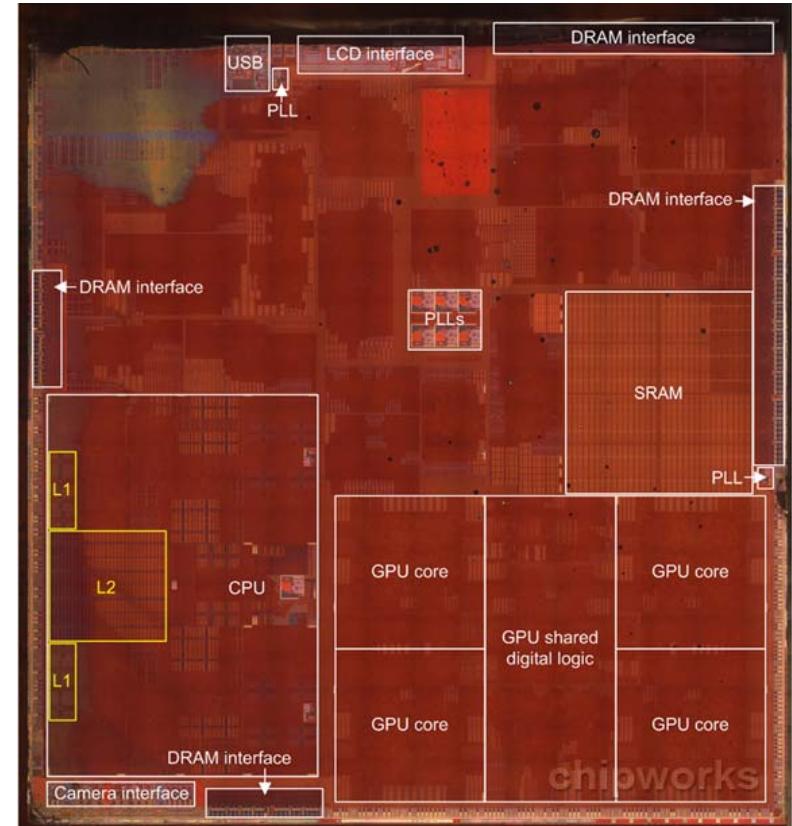
Modern Processors

3rd Generation Intel® Core™ Processor:
22nm Process

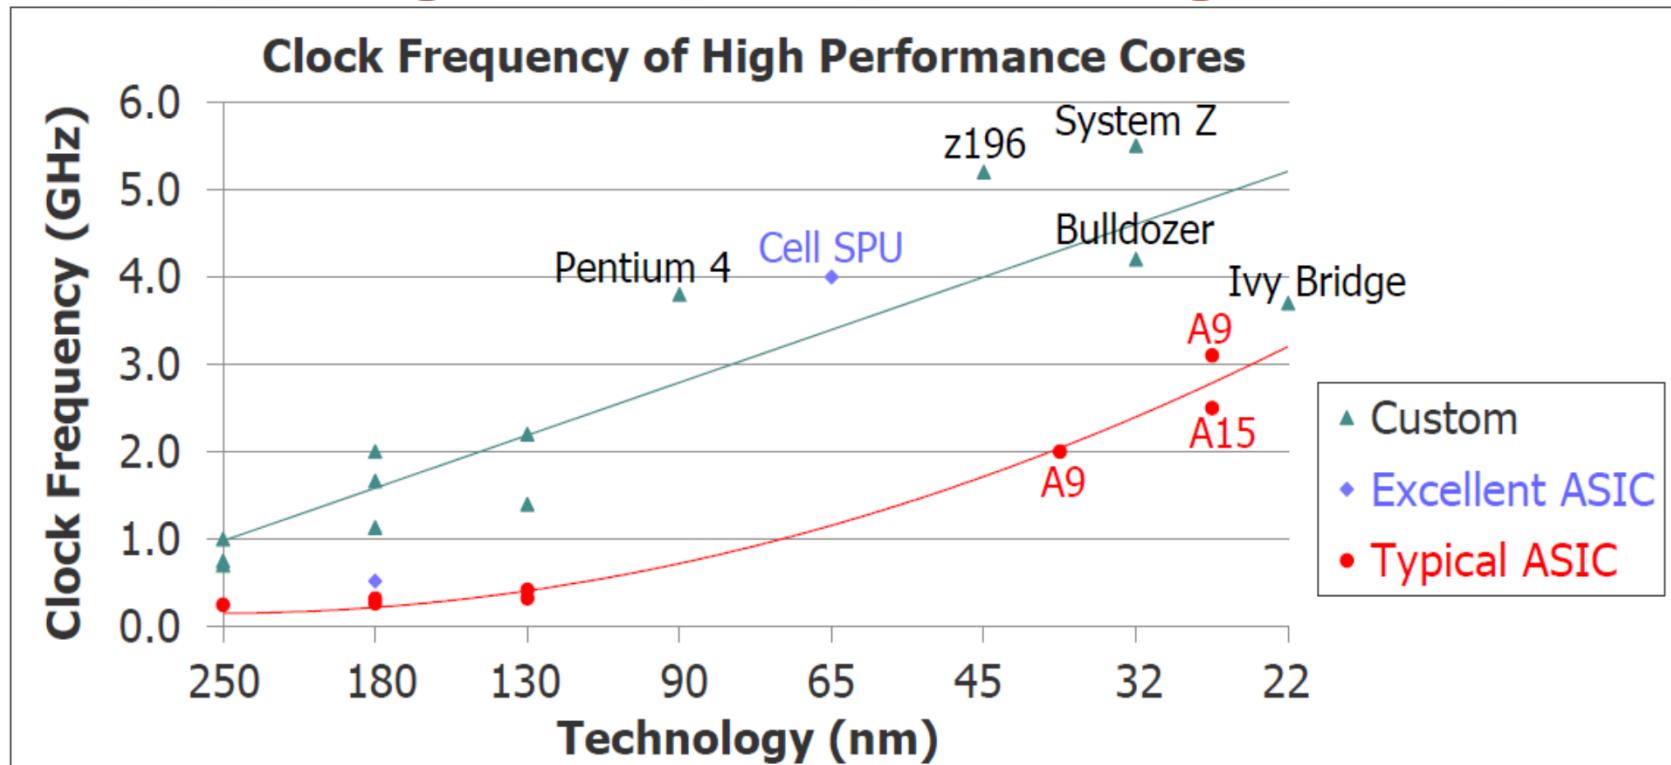


New architecture with shared cache delivering more performance and energy efficiency

Quad Core die with Intel® HD Graphics 4000 shown above
Transistor count: 1.4Billion Die size: 160mm²
** Cache is shared across all 4 cores and processor graphics.



Clocking in Modern Designs



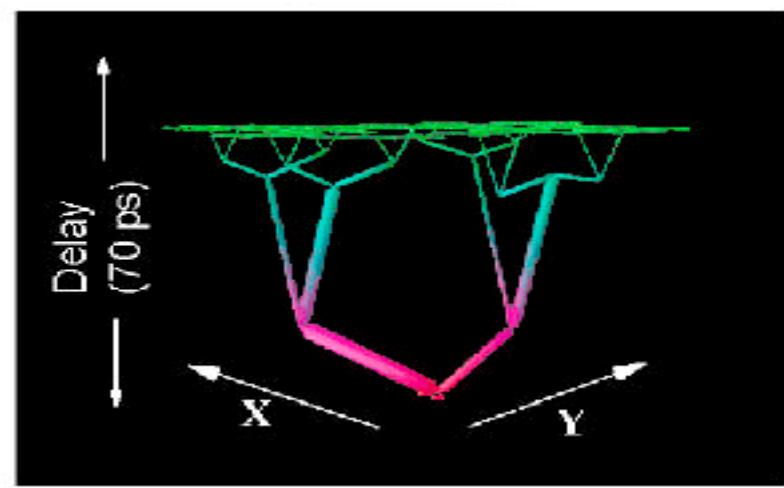
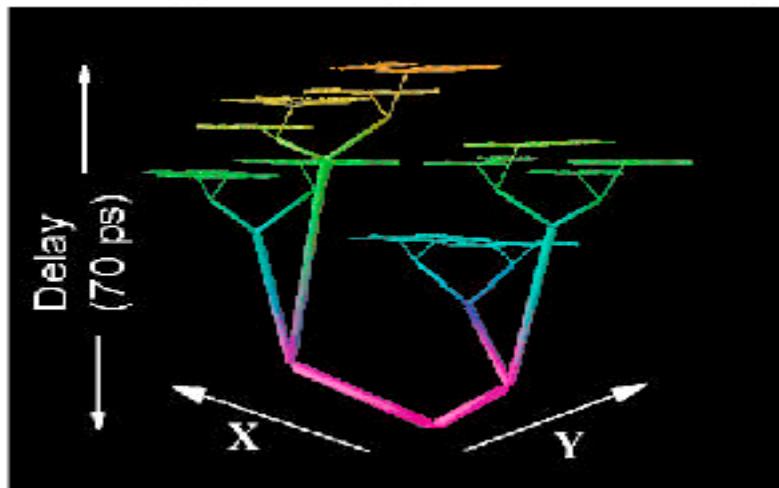
David Chinnery Mentor Graphics

- ❑ Custom designs were 3 to 8 × faster than ASICs
- ❑ Performance gap is below 2 × today, custom limited by long design time

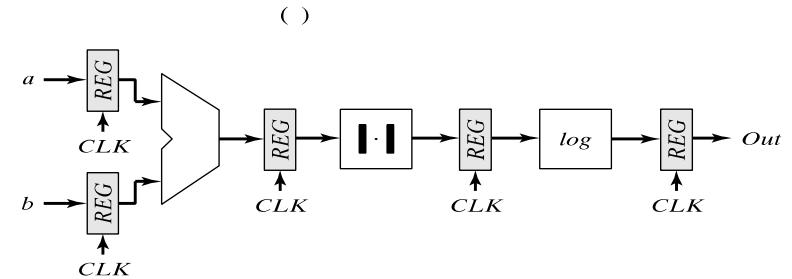
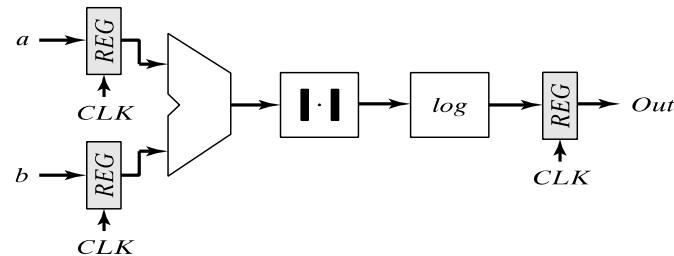
Clock Animations

□ By Phillip Restle (IBM)

<http://www.research.ibm.com/people/r/restle/resonate.html>



Pipelining

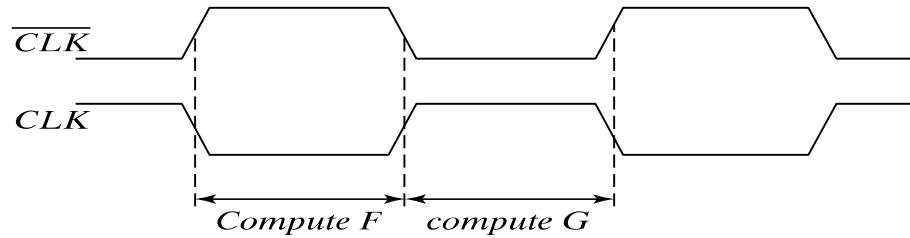
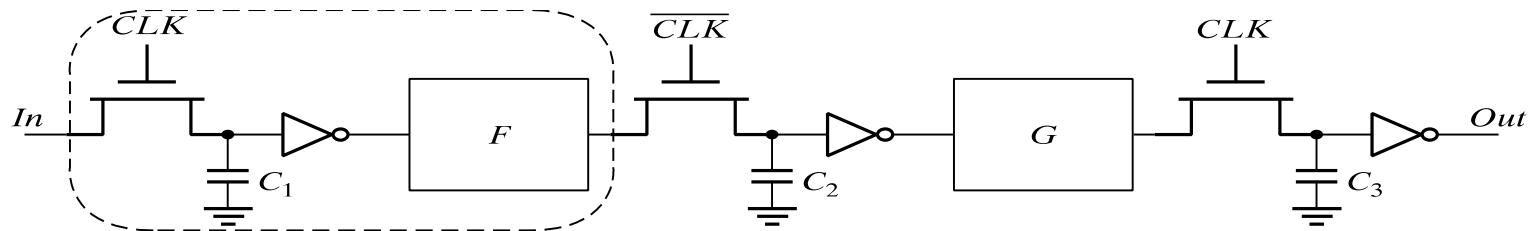


Reference

Clock Period	Adder	Absolute Value	Logarithm
1	$a_1 + b_1$		
2	$a_2 + b_2$	$ a_1 + b_1 $	
3	$a_3 + b_3$	$ a_2 + b_2 $	$\log(a_1 + b_1)$
4	$a_4 + b_4$	$ a_3 + b_3 $	$\log(a_2 + b_2)$
5	$a_5 + b_5$	$ a_4 + b_4 $	$\log(a_3 + b_3)$

Pipelined

Latch-Based Clocking

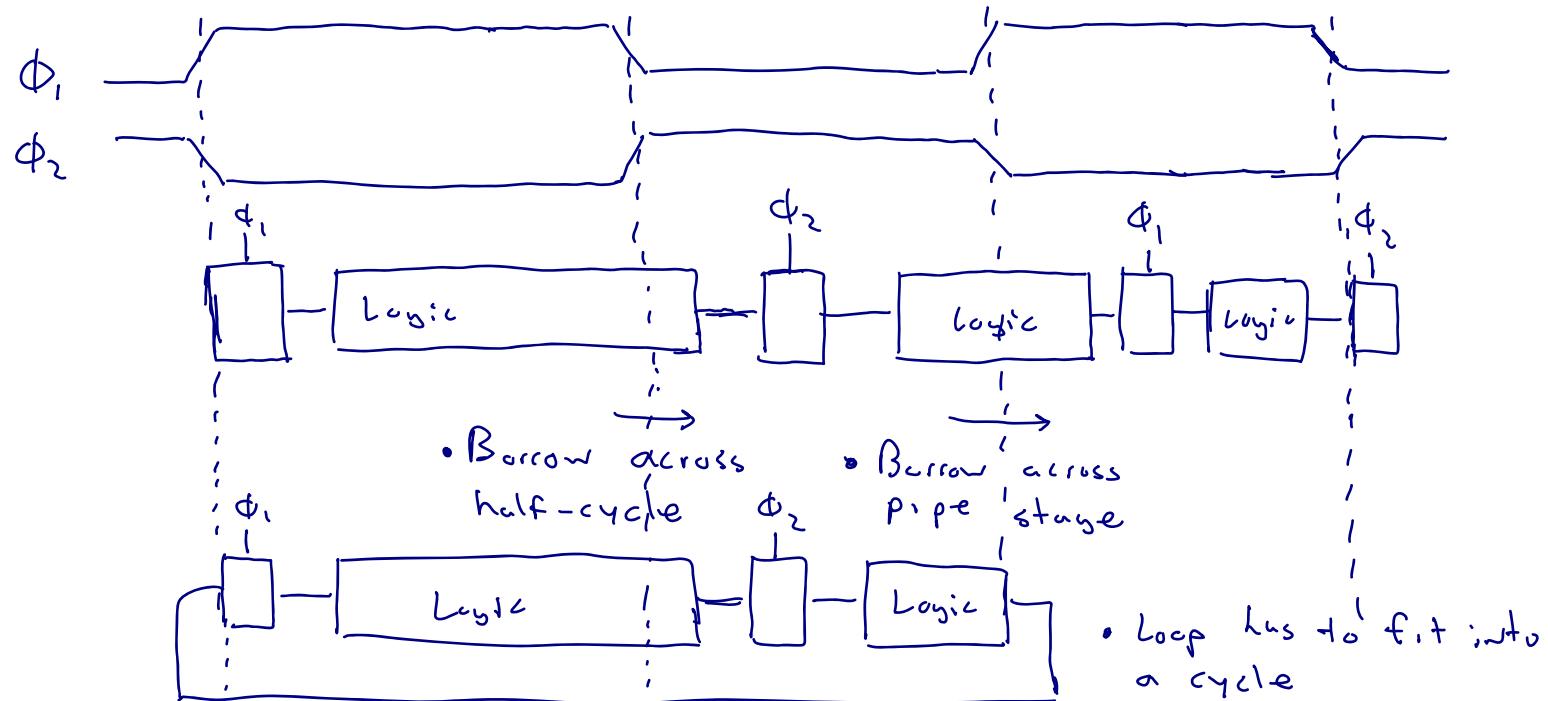


(Domino logic almost always uses latch-based clocking)

Latch vs. Flip-flop

- In a flip-flop based system:
 - Data launches on one rising edge
 - And must arrive before next rising edge
 - If data arrives late, system fails
 - If it arrives early, wasting time
 - Flip-flops have hard edges
- In a latch-based system:
 - Data can pass through latch while it is transparent
 - Long cycle of logic can borrow time into next cycle
 - As long as each loop finished in one cycle

Time Borrowing Example



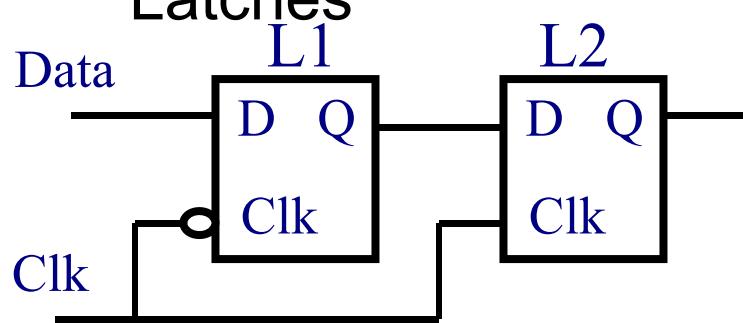
Latch vs. Flip-flop Summary

- Flip-flops generally easier to use
 - Most digital ASICs designed with register-based timing
- But, latches (both pulsed and level-sensitive) allow more flexibility
 - And hence can potentially achieve higher performance
 - Latches can also be made more tolerant of clock un-certainty
 - More in “Advanced Digital Logic” Class

Pulse-Triggered Latches

Ways to design an edge-triggered sequential cell:

Master-Slave
Latches



Pulse-Triggered
Latch

